

Saeed Safari,
Advanced Computer Architecture

Introduction

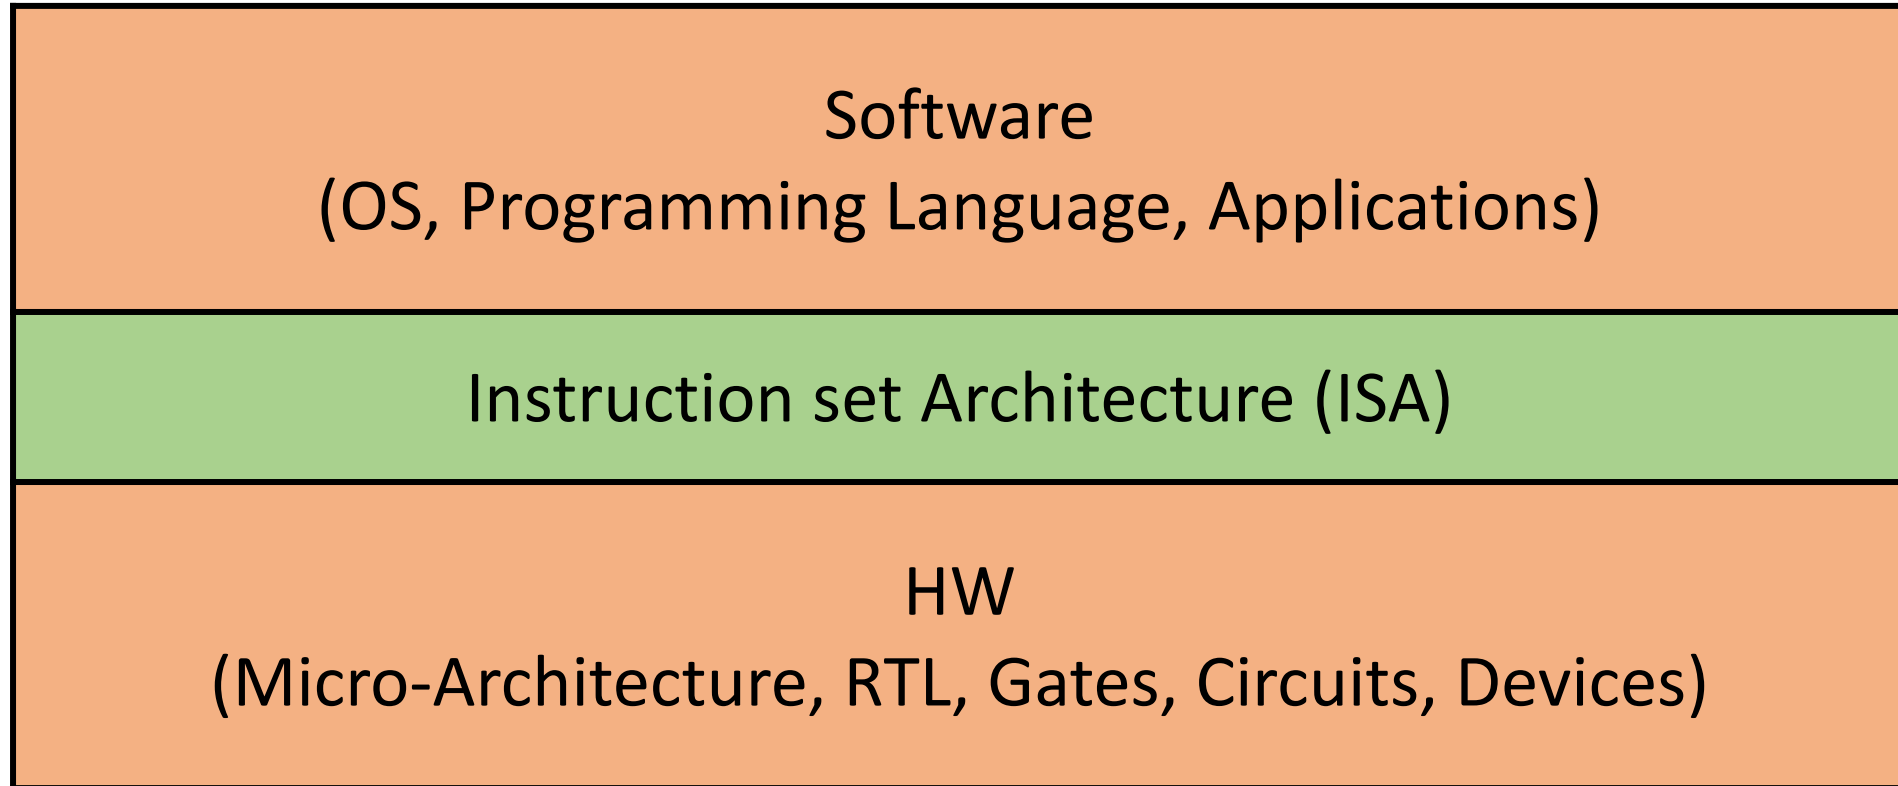
*ECE School,
Faculty of Engineering,
University of Tehran*

Outline

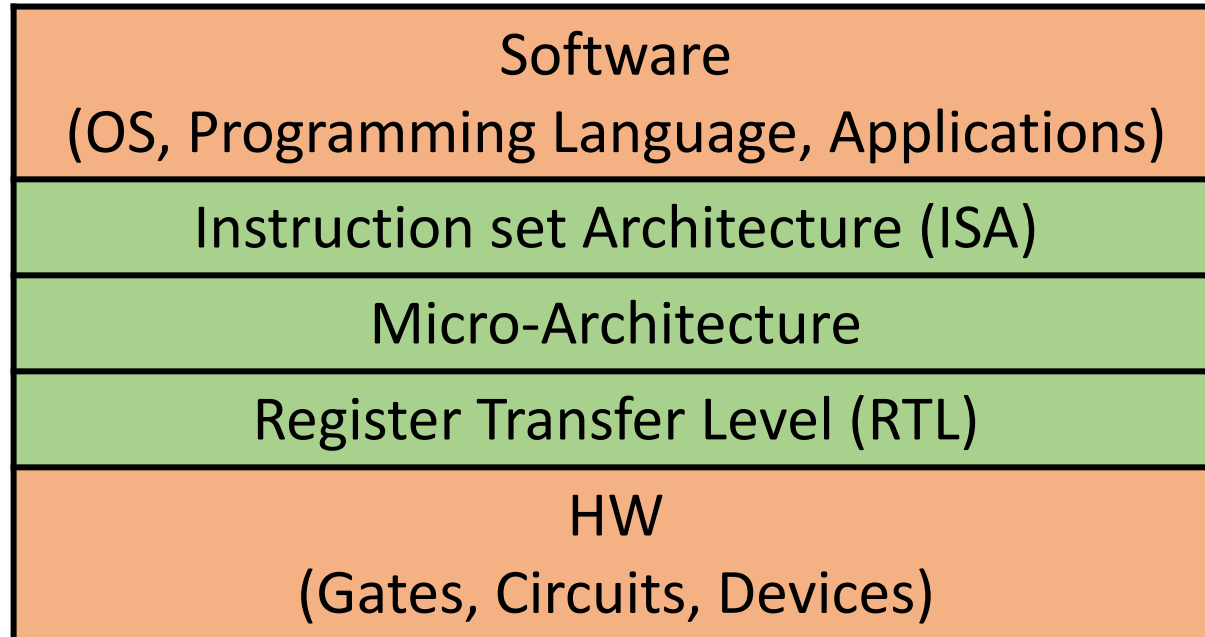
- Architecture
 - CPU classes
 - Addressing modes
 - Operands type and size
 - Operations
 - Instruction encoding
- Micro-Architecture



Computing Systems Abstraction Levels




Computing Systems Abstraction Levels



**Computer Architecture
and
Advanced Computer Architecture**

Single Processor Performance



Architecture vs. Micro-Architecture

- Architecture / Instruction Set Architecture (ISA)
 - Programmer visible state (Registers, Memory)
 - Data types (Integer, FP)
 - Addressing modes
 - Operations
 - Instruction encoding
 - ...



Architecture vs. Micro-Architecture

- Micro-Architecture / Organization
 - Tradeoffs on ISA implementation (Speed, cost, power, ...)
 - **Examples:**
 - Pipeline depth,
 - Number of pipelines,
 - ALU width,
 - Cache size,
 - Execution ordering,
 - ...



Same Architecture Different Micro-Architecture

● Intel Atom

- X86 Instruction Set
- Single Core
- 2 W
- Decode 2 Instructions / Cycle / Core
- 32 KB L1 I Cache, 32 KB L1 D Cache
- 512 KB L2 Cache
- In-Order
- 47,000,000 Tr.
- 1.6 GHz




Image Credit: Intel

● AMD Phenom X4

- X86 Instruction Set
- Quad Core
- 125 W
- Decode 3 Instructions / Cycle / Core
- 64 KB L1 I Cache, 64 KB L1 D Cache
- 512 KB L2 Cache
- Out-Of-Order
- 758,000,000 Tr.
- 2.6 GHz




Image Credit: AMD

Different Architecture Different Micro-Architecture

● IBM POWER7

- Power Instruction Set
- Eight Core
- 200 W
- Decode 6 Instructions / Cycle / Core
- 32 KB L1 I Cache, 32 KB L1 D Cache
- 256 KB L2 Cache
- Out-Of-Order
- 1,200,000,000 Tr.
- 4.25 GHz




Image Credit: IBM

● AMD Phenom X4

- X86 Instruction Set
- Quad Core
- 125 W
- Decode 3 Instructions / Cycle / Core
- 64 KB L1 I Cache, 64 KB L1 D Cache
- 512 KB L2 Cache
- Out-Of-Order
- 758,000,000 Tr.
- 2.6 GHz




Image Credit: AMD

Architecture

- **Complete:** should be able to implement all user program
 - **Concise:** limited size of the instruction set
 - **Generic:** instructions should capture the common case
 - **Simple:** instructions should be simple
-
- It is necessary to be complete, whereas the rest of the properties are desirable

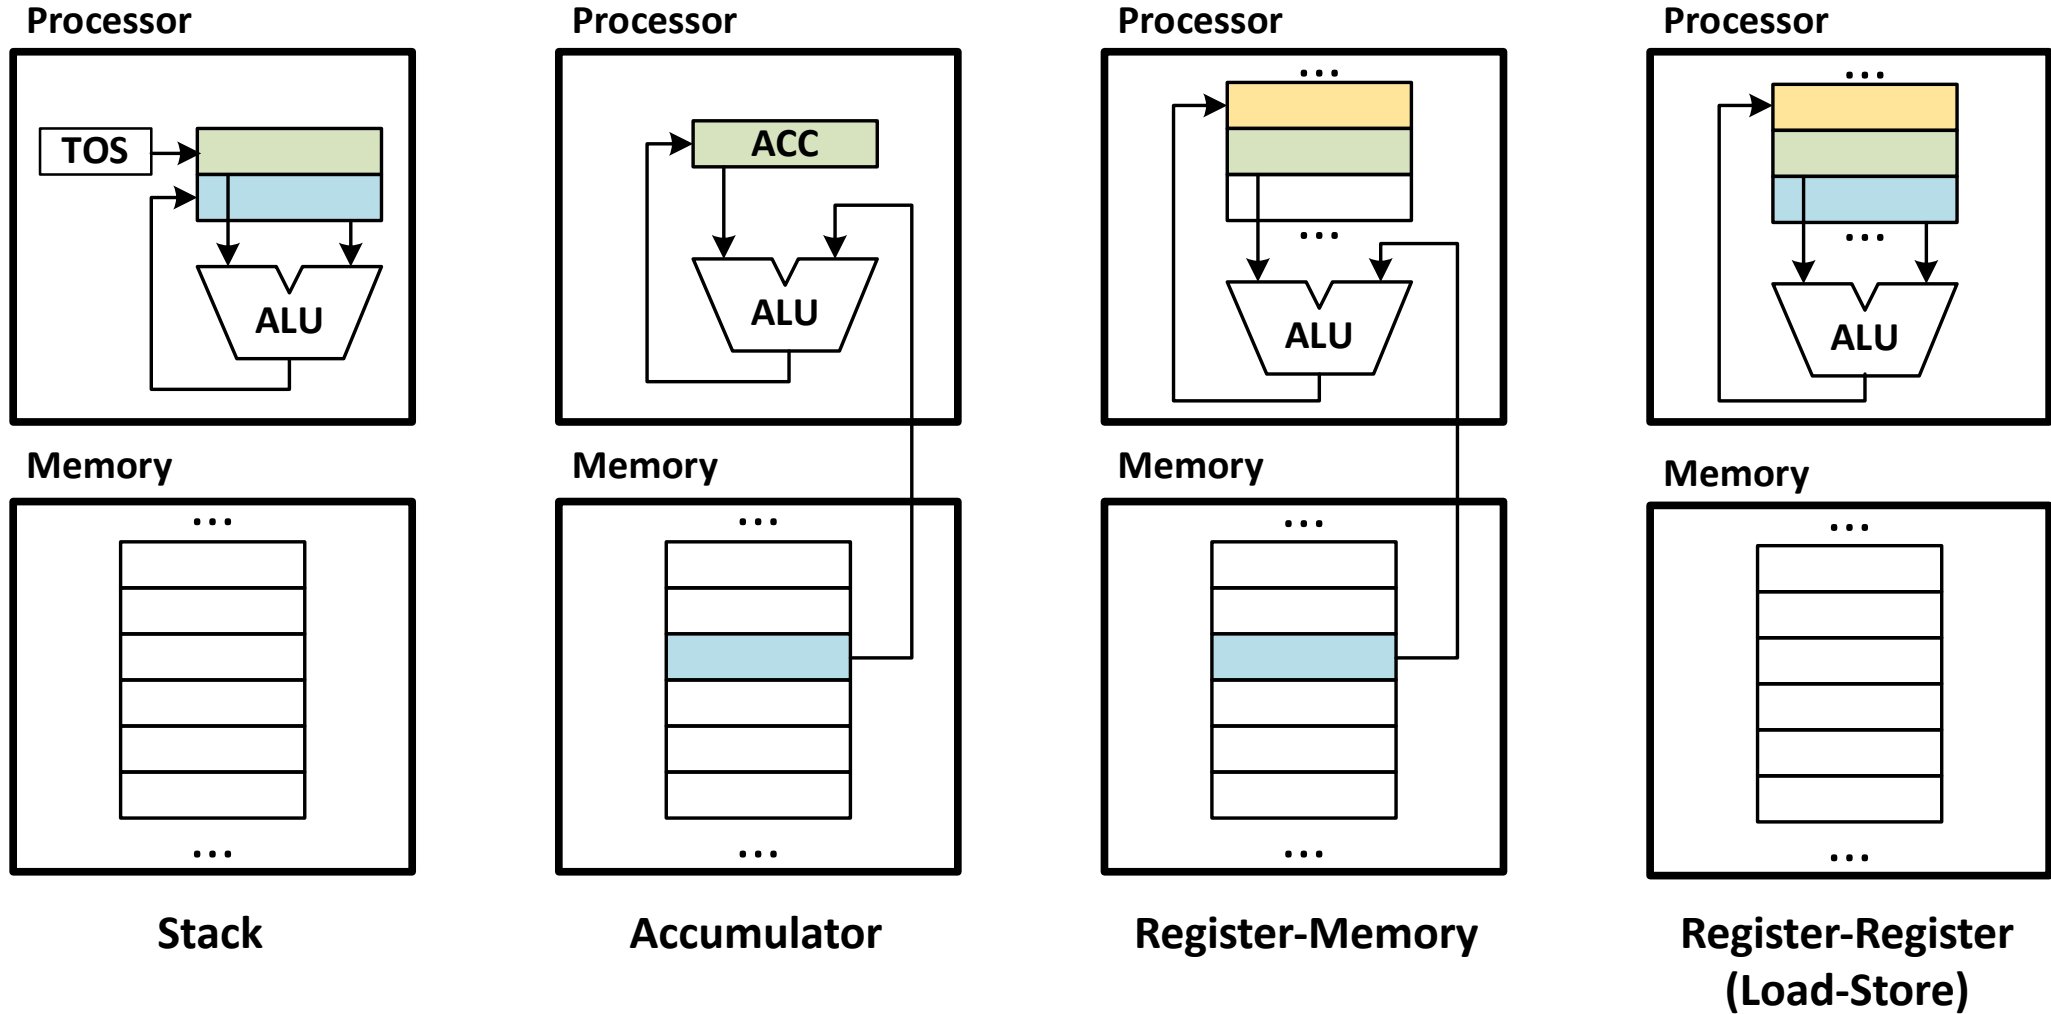


Architecture


- Machine model
- Memory addressing
- Type and size of operands
- Operations
- Encoding



Machine Models




Machine Models



$C = A + B;$

PUSH A
PUSH B
ADD
POP C


Machine Models



$C = A + B;$

LOAD A
ADD B
STORE C


Machine Models



$C = A + B;$

LOAD R1, A
ADD R3, R1, B
STORE R3, C

Machine Models



$C = A + B;$

LOAD R1, A

LOAD R2, B

ADD R3, R1, R2


STORE R3, C

Machine Models


Architecture Type	Example
Stack	HP 3000, Intel X87 FPU
Load-store	ARM, MIPS, PowerPC, SPARC, RISC-V
Register-memory	IBM 360/370, Intel X86, Motorola 68000



Stack-Based Processor





$$y = (a + b * c) / (a + d * c - e)$$



● Reverse Polish

$a \ b \ c \ * \ + \ a \ d \ c \ * \ + \ e \ - \ /$

Stack-Based Processor



a b c * + a d c * + e - /

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack



Stack-Based Processor

PUSH **a**

PUSH b

PUSH c

MUL

ADD

PUSH a

PUSH d

PUSH c

MUL

ADD

PUSH e

SUB

DIV

POP y

Stack

a



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

b
a



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

c
b
a



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

$b * c$
a



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

$a+b*c$



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

a
$a+b*c$



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

d
a
$a+b*c$



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

c
d
a
$a+b*c$



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

$d * c$
a
$a + b * c$



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

$a+d*c$
$a+b*c$



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

e
$a+d*c$
$a+b*c$



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

$a+d*c-e$
$a+b*c$



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV
POP y

Stack

$(a+b*c)/(a+d*c-e)$



Stack-Based Processor

PUSH a
PUSH b
PUSH c
MUL
ADD
PUSH a
PUSH d
PUSH c
MUL
ADD
PUSH e
SUB
DIV

POP **y**

Stack



Stack Hardware Organization

- Stack is part of the processor state
 - Stack must be bounded and small
 - Number of registers
 - Smaller than the main memory
- Conceptually stack is unbounded
 - A part is in the processor state
 - The rest is in the main memory



Stack Hardware Organization

- Each PUSH (or POP) operation needs 1 memory reference
- Assumption
 - Stack size = 2 (4)
 - The 2 (4) top elements stored in registers
 - The rest is in main memory



Stack Hardware Organization

Program	Stack (Size = 2)	Memory References	
PUSH a	R0	a	
PUSH b	R0 R1	b, a	
PUSH c	R0 R1 R2	c, b, ss(a)	4 implicit stores
MUL	R0 R1	b*c, sf(a)	4 implicit fetches
ADD	R0	a+b*c	
PUSH a	R0 R1	a, a+b*c	
PUSH d	R0 R1 R2	d, a, ss(a+b*c)	
PUSH c	R0 R1 R2 R3	c, d, ss(a)	
MUL	R0 R1 R2	d*c, sf(a)	
ADD	R0 R1	a+d*c, sf(a+b*c)	
PUSH e	R0 R1 R2	e, a+d*c, ss(a+b*c)	
SUB	R0 R1	a+d*c-e, sf(a+b*c)	
DIV	R0	(a+b*c) / (a+d*c-e)	

Stack Hardware Organization

Program		Stack (Size = 4)
PUSH	a	R0
PUSH	b	R0 R1
PUSH	c	R0 R1 R2
MUL		R0 R1
ADD		R0
PUSH	a	R0 R1
PUSH	d	R0 R1 R2
PUSH	c	R0 R1 R2 R3
MUL		R0 R1 R2
ADD		R0 R1
PUSH	e	R0 R1 R2
SUB		R0 R1
DIV		R0

a and c are loaded twice



Memory Addressing

- How many address lines?
 - 32-bit
 - 64-bit
- Byte / Word addressable?



Memory Addressing

- Little / Big Endian?
 - Little endian



- Big endian



Memory Addressing

● Aligned / Unaligned?

$$Adr \bmod size = \begin{cases} 0 & \text{aligned} \\ 1 & \text{unaligned} \end{cases}$$

Width of objects	0	1	2	3	4	5	6	7
1 byte (byte)	Aligned	Aligned	Aligned	Aligned	Aligned	Aligned	Aligned	Aligned
2 bytes (half word)	Aligned		Aligned		Aligned		Aligned	
2 bytes (half word)		Unaligned		Unaligned		Unaligned		Unaligned
4 bytes (word)	Aligned				v			
4 bytes (word)		Unaligned				Unaligned		
4 bytes (word)			Unaligned				Unaligned	
4 bytes (word)			Unaligned					Unaligned

Memory Addressing

Addressing mode	Example	Meaning
Register	Add R1, R2	$Regs[R1] \leftarrow Regs[R1] + Regs[R2]$
Immediate	Add R1, 3	$Regs[R1] \leftarrow Regs[R1] + 3$
Displacement	Add R1, 100(R2)	$Regs[R1] \leftarrow Regs[R1] + Mem[100 + Regs[R2]]$
Register Indirect	Add R1, (R2)	$Regs[R1] \leftarrow Regs[R1] + Mem[Regs[R2]]$
Indexed	Add R1, (R2+R3)	$Regs[R1] \leftarrow Regs[R1] + Mem[Regs[R2] + Regs[R3]]$
Direct (Absolute)	Add R1, (100)	$Regs[R1] \leftarrow Regs[R1] + Mem[100]$
Memory Indirect	Add R1, @(R2)	$Regs[R1] \leftarrow Regs[R1] + Mem[Mem[Regs[R2]]]$



Memory Addressing

Addressing mode	Example	Meaning
Autoincrement	Add R1, (R2)+	$Regs[R1] \leftarrow Regs[R1] + Mem[Regs[R2]]$ $Regs[R2] \leftarrow Regs[R2] + d$
Autodecrement	Add R1, -(R2)	$Regs[R2] \leftarrow Regs[R2] - d$ $Regs[R1] \leftarrow Regs[R1] + Mem[Regs[R2]]$
Scaled	Add R1, 100(R2)[R3]	$Regs[R1] \leftarrow Regs[R1] +$ $Mem[100 + Regs[R2] + Regs[R3] * d]$



Type and Size of Operands

- Types:
 - Integer (Signed / Unsigned)
 - Binary Coded Decimal (BCD)
 - Floating Point
 - IEEE-754
 - Intel Extended Precision (80-bit)
 - Packed Vector Data



Type and Size of Operands

- Sizes:
 - Integer (8-bit, 16-bit, 32-bit, 64-bit)
 - Floating Point
 - Single Precision (SP: 32-bit)
 - Double Precision (DP: 64-bit)
 - Intel Extended Precision (80-bit)



Operations

- Data transfer:
 - LD, ST, MFC1, MTC1
- ALU operation:
 - ADD, SUB, AND, OR, XOR, MUL, DIV, SLT, LUI
- Floating Point:
 - ADD.S, SUB.D, MUL.D, CVT.S.W, CVT.W.S
- Control Flow:
 - J, JR, BEQ, JAL



Operations

- Multimedia (SIMD):
 - ADD.PS, SUB.PD, MUL.PS
- String
 - REP MOVSB (x86)



ISA Encoding


- Fixed Width:
 - Every instruction has the same width
 - Easy to decode
 - RISC architectures :
 - MIPS, PowerPC, SPARC, ARM, RISC-V, ...
 - Example: MIPS, every instruction has 4-bytes

Operation	Address field 1	Address field 2	Address field 3
-----------	--------------------	--------------------	--------------------



ISA Encoding

- Variable Length:
 - Instruction can have different width
 - Less space in cache / memory
 - CISC architectures:
 - IBM 360, Intel x86, Motorola 68K, VAX, ...
 - Example: X86, instructions are 1-byte up to 17-bytes



ISA Encoding

- Mostly Fixed:
 - Instruction has two different widths
 - Example:
 - MIPS16: two types of instructions: 2-bytes and 4-bytes
 - THUMBS (ARM16): two types of instructions: 2-bytes and 4-bytes



ISA Encoding

- Very Long Instruction Word (VLIW):
 - Multiple instructions in a fixed length bundle
 - Example:
 - Multiflow
 - HP/ST* Lx
 - TI C6000

* Semiconductor Technology



Real World Instruction Sets

Arch	Type	# Opnd	# Mem	Data Size	# Regs	Adr Size	Use
Alpha	Reg-Reg	3	0	64-bit	32	64-bit	Workstation
ARM	Reg-Reg	3	0	32/64-bit	16	32/64-bit	Cell Phones, Embedded
MIPS	Reg-Reg	3	0	32/64-bit	32	32/64-bit	Workstation, Embedded
SPARC	Reg-Reg	3	0	32/64-bit	24-32	32/64-bit	Workstation
TI C6000	Reg-Reg	3	0	32-bit	32	32-bit	DSP
IBM 360	Reg-Mem	2	1	32-bit	16	24/31/64	Mainframes
x86	Reg-Mem	2	1	8/16/32/64-bit	4/8/24	16/32/64	Personal Computers
VAX	Mem-Mem	3	3	32-bit	16	32-bit	Minicomputer
Mot. 6800	Accumulator	1	1/2	8-bit	0	16-bit	Microcontroller



Summary

- Topics covered
 - Architecture vs Micro-architecture
 - Architecture
 - Machine model
 - Memory addressing
 - Data types
 - Instructions
 - Encoding
 - Micro-architecture

