

Saeed Safari,
Advanced Computer Architecture

Pipeline – Basic Concepts

*ECE School,
Faculty of Engineering,
University of Tehran*

Outline

- MIPS implementation
 - Single-cycle
 - Multi-cycle
 - Pipeline
- Hazard
 - Structural hazard
 - Data hazard
 - Control hazard



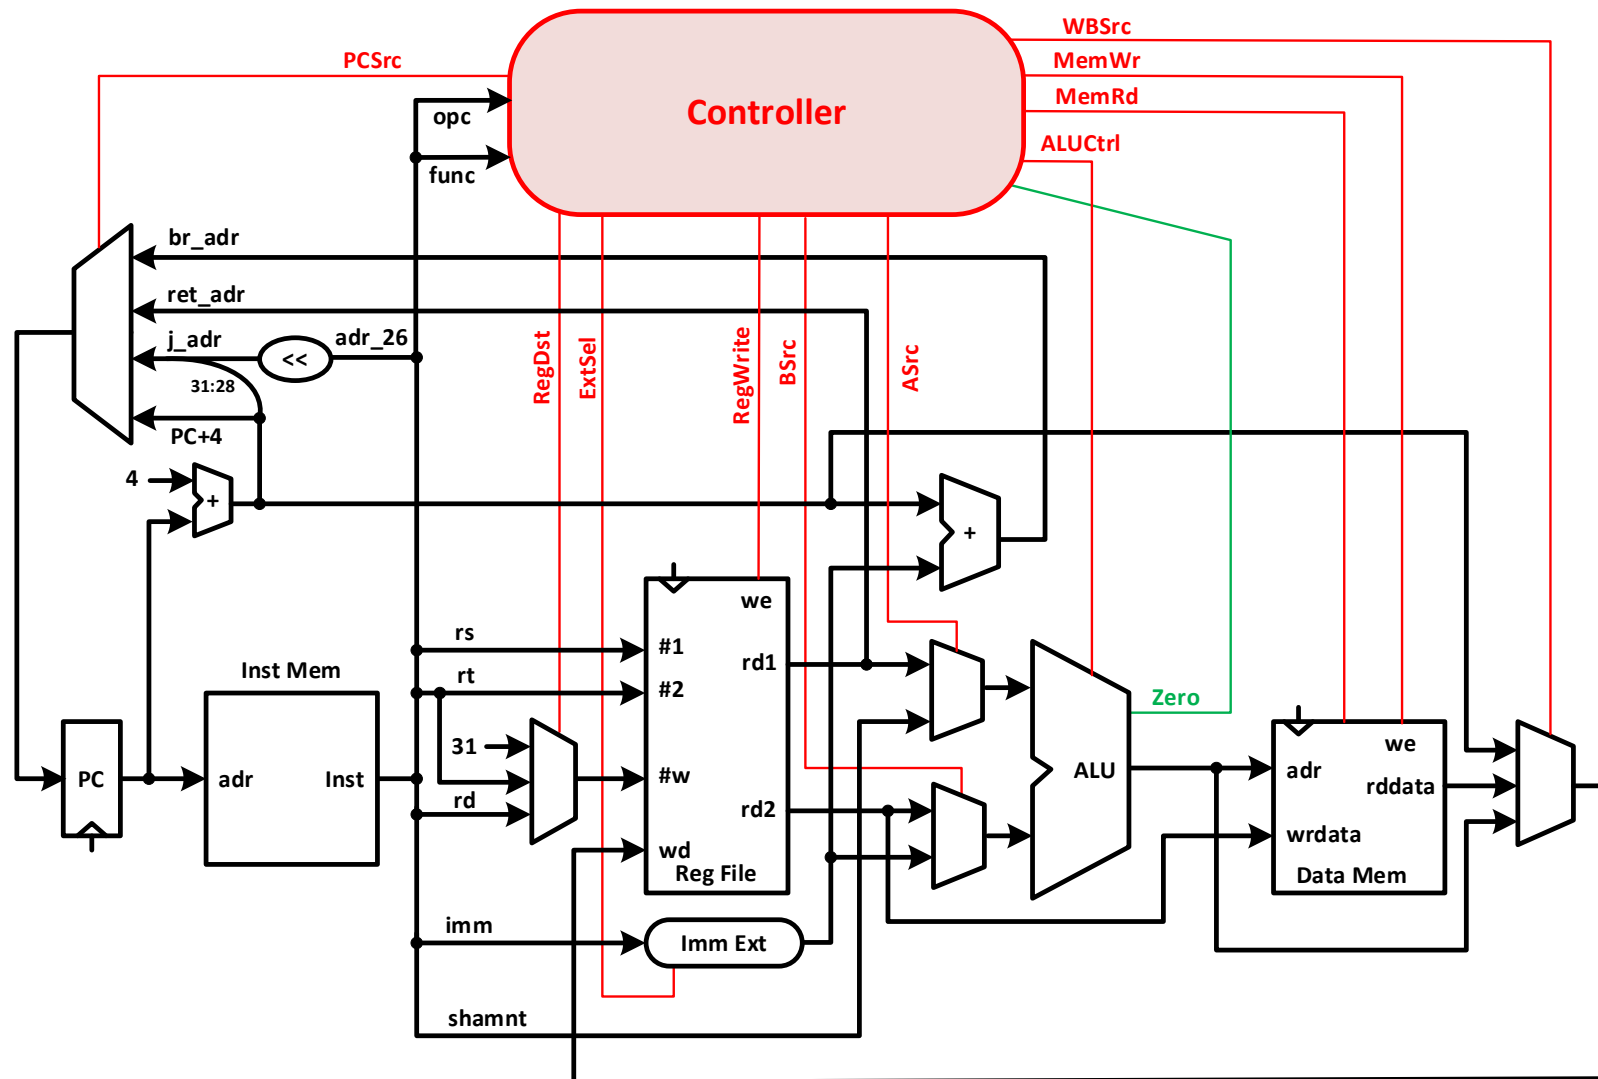
MIPS ISA – An Overview

| | |
|-------------------------|---|
| Registers | 32 32-bit integer registers (R0 – R31), 32 32-bit floating point registers (F0 – F31), 2 32-bit integer register (hi, lo) |
| Addressing Modes | Register, Immediate, PC Relative, Base |

| | |
|----------------|---|
| R-Type | add, addu, sub, subu, slt, sltu, and, or, xor, nor, sll, srl, sla, sllv, arlv, slav, jr, jalr, mult, multu, div, divu, mfhi, mthi, mflo, mtlo |
| I-Type | addi, addiu, slti, sltiu, andi, ori, xori, lui, lw, sw, beq, bne |
| J-Type | j, jal |
| FR-Type | add.s, sub.s, mul.s, div.s, abs.s, neg.s, c.eq.s, c.lt.s, c.le.s add.d, sub.d, mul.d, div.d, abs.d, neg.d, c.eq.d, c.lt.d, c.le.d |
| FI-Type | l.s, s.s, l.d, s.d, bc1t, bc1f |



MIPS – Single Cycle Implementation



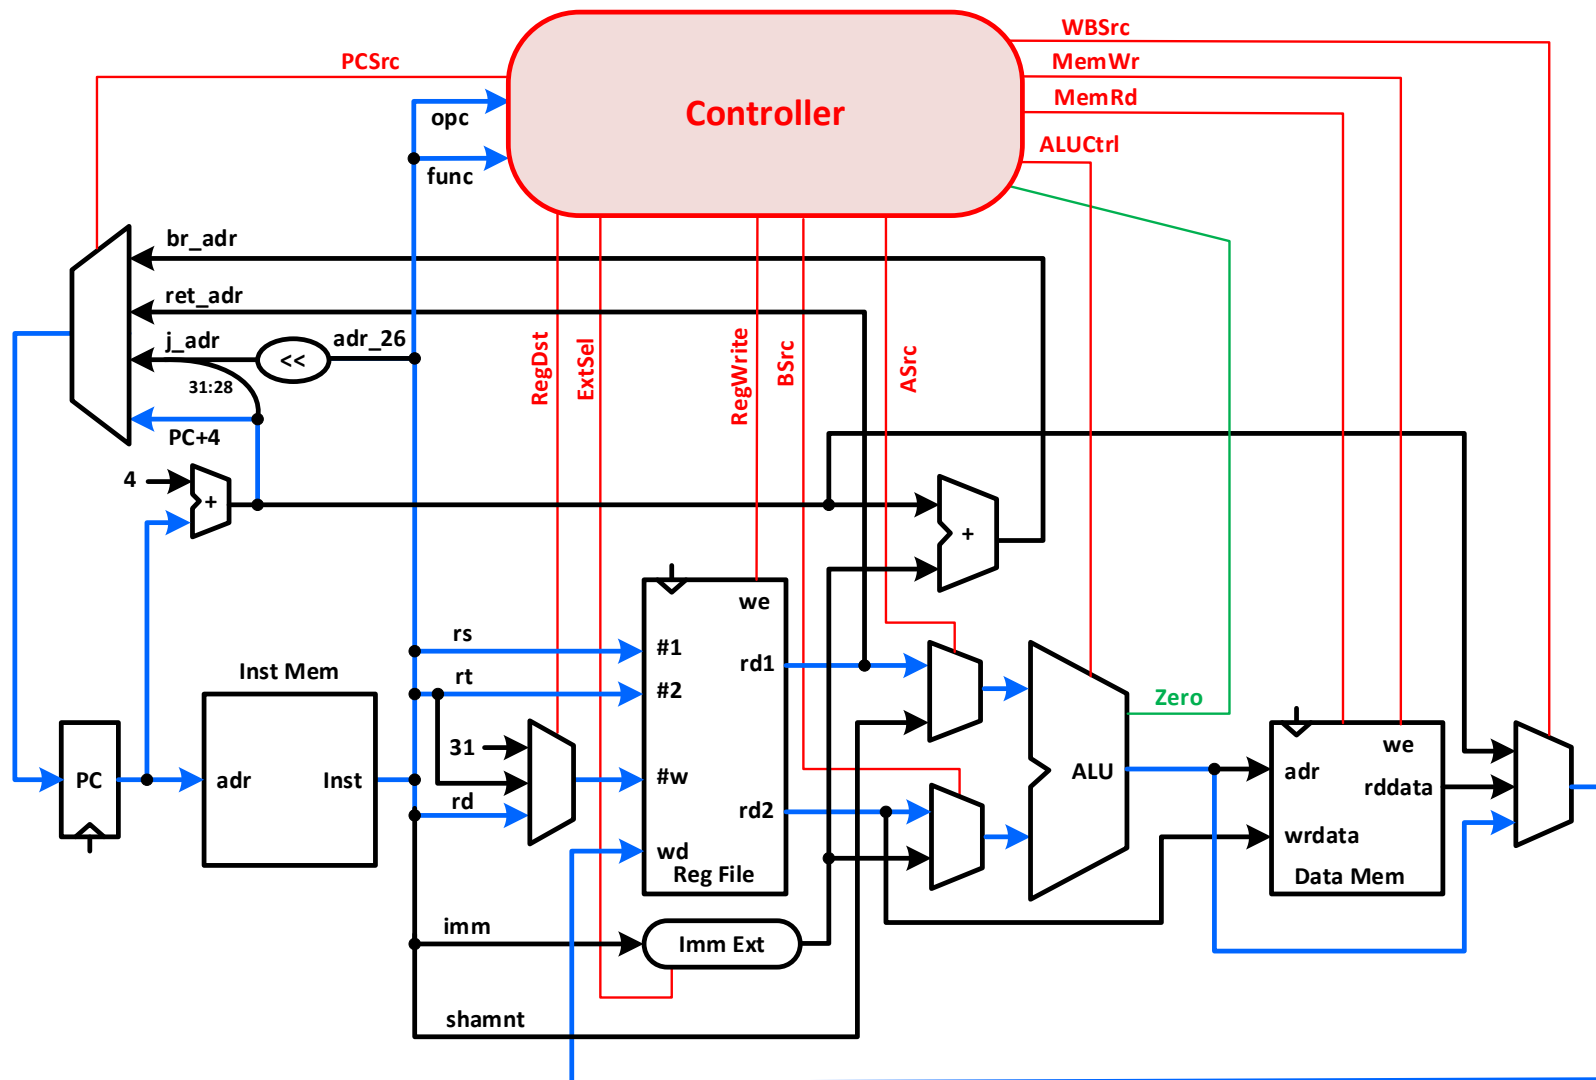
MIPS – Single Cycle Implementation

- Key features:
 - Two different parts
 - Data-path
 - Controller
 - Controller is implemented as a combinational circuit
 - Each instruction needs one clock cycle to execute (single-cycle implementation)



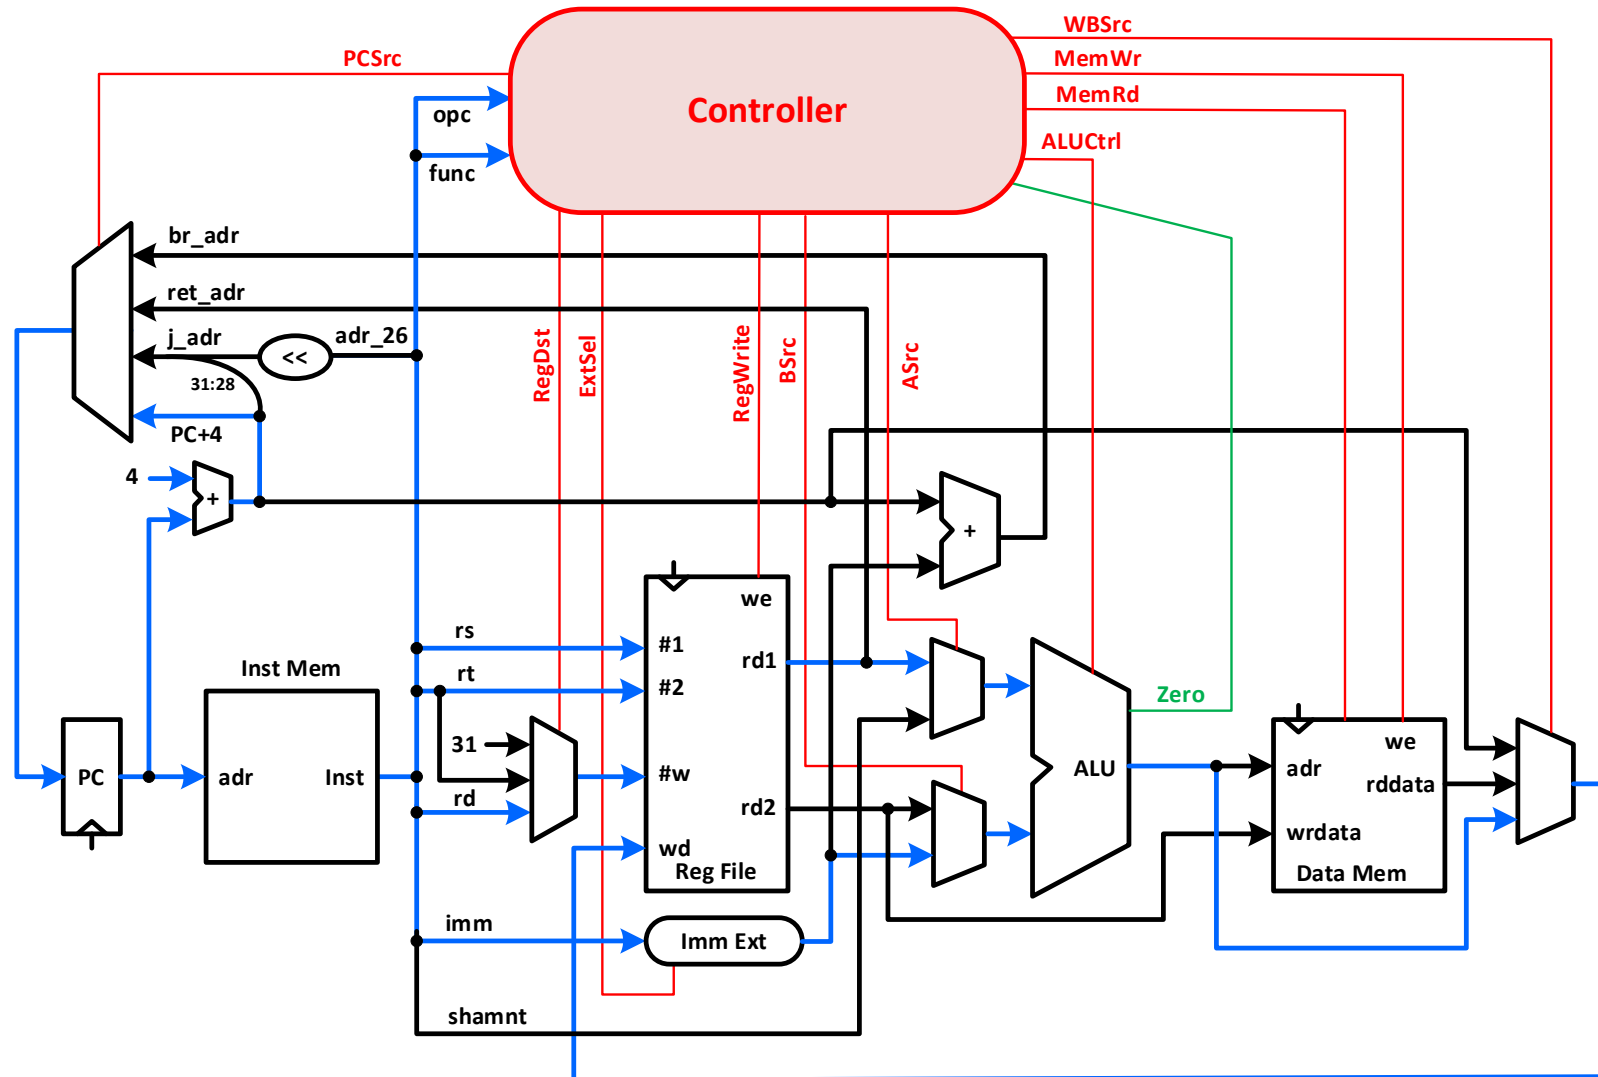
Arithmetic–Logic Instructions

add,
addu,
sub,
subu,
slt,
sltu,
and,
or,
xor,
nor

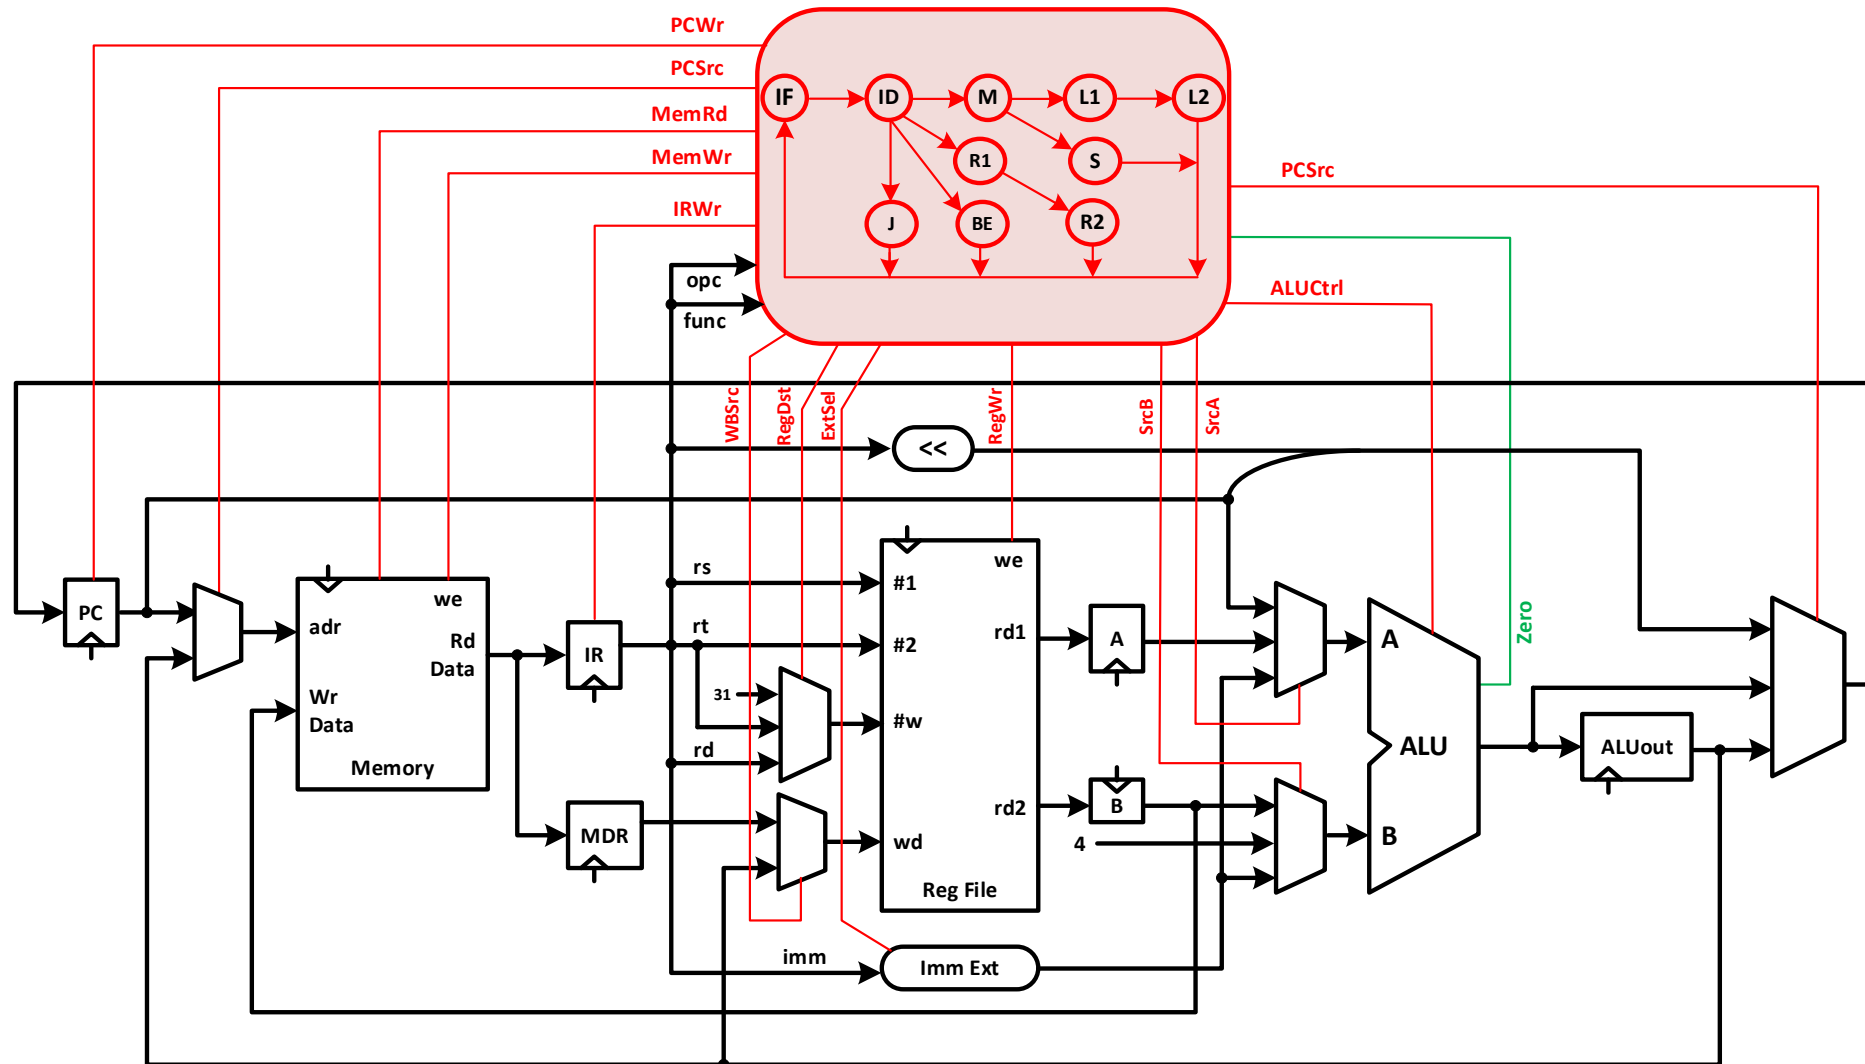


Immediate Instructions

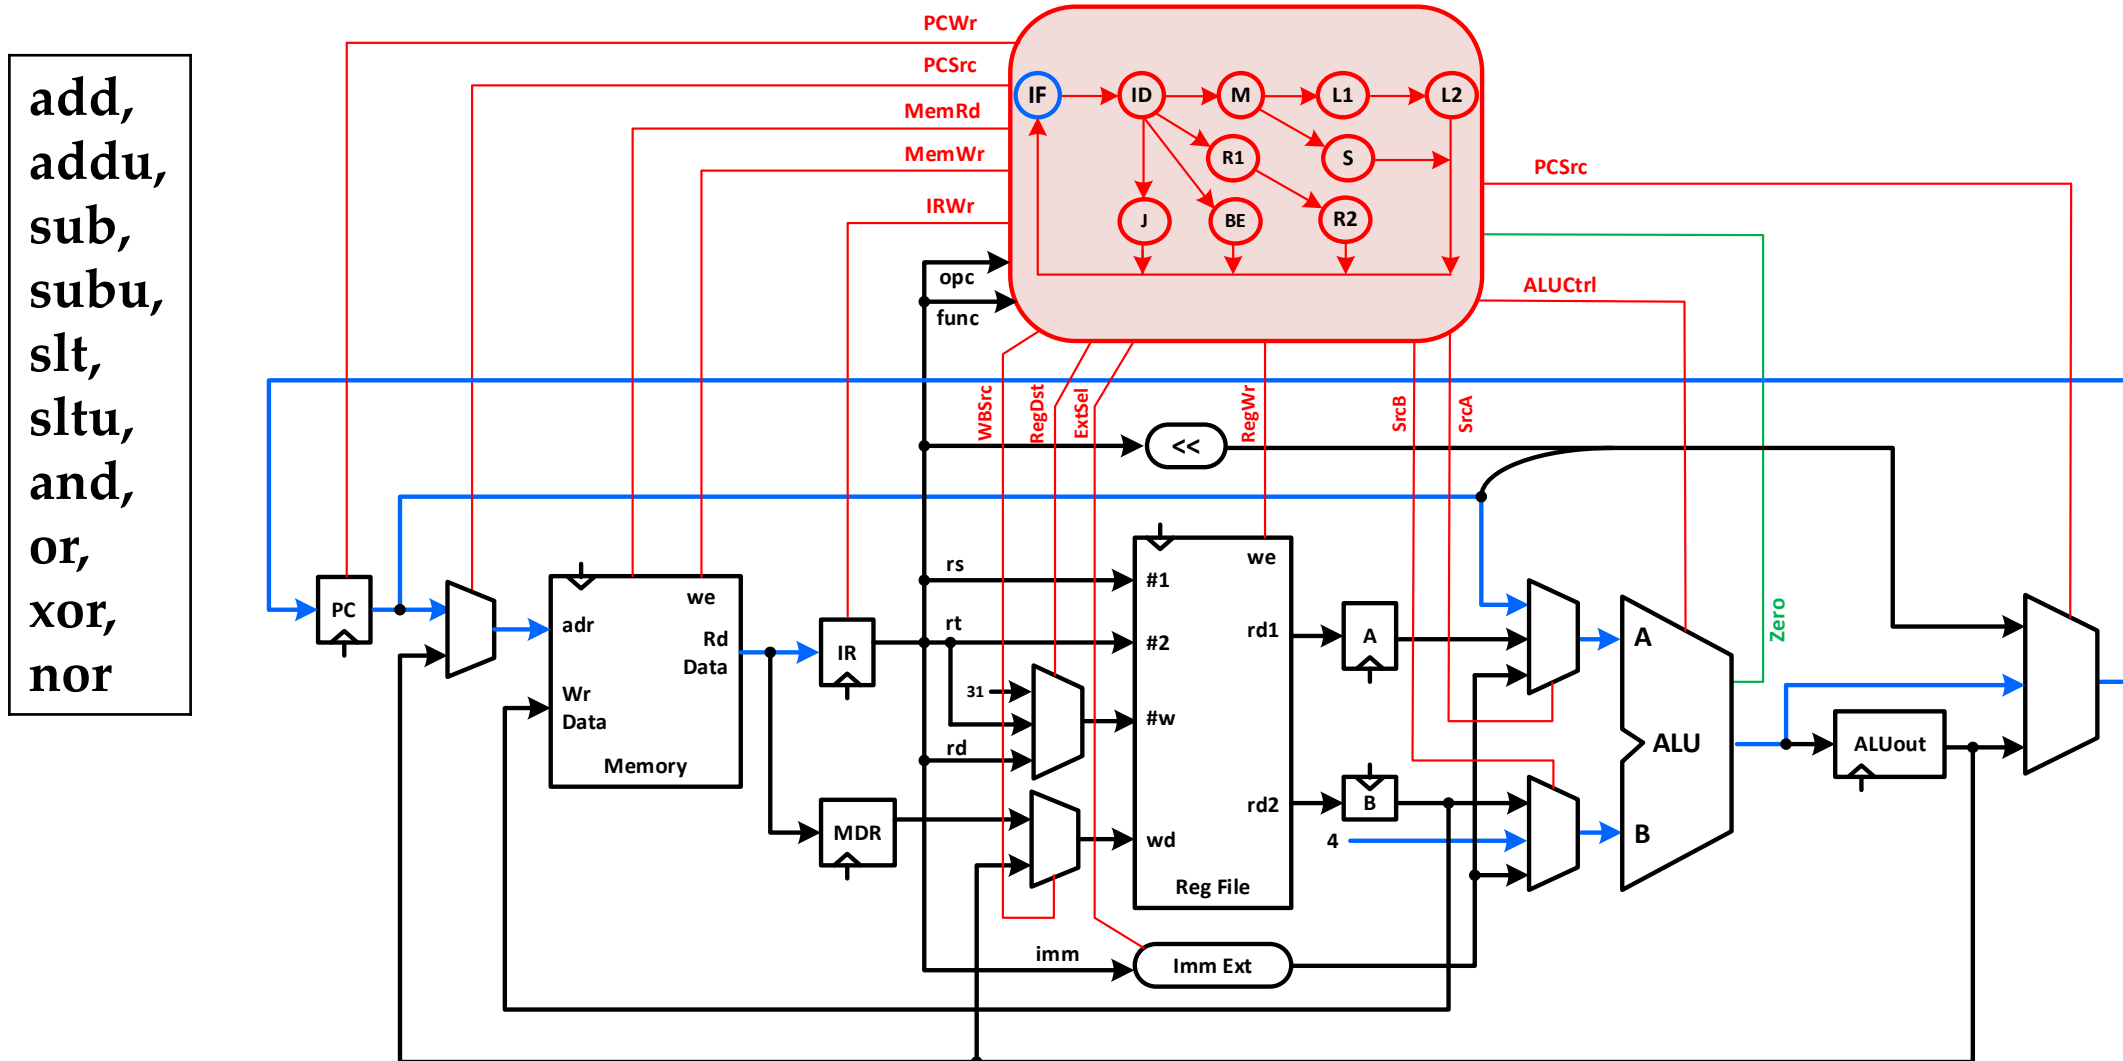
addi,
addiu,
slti,
sltiu,
andi,
ori,
xori,
lui



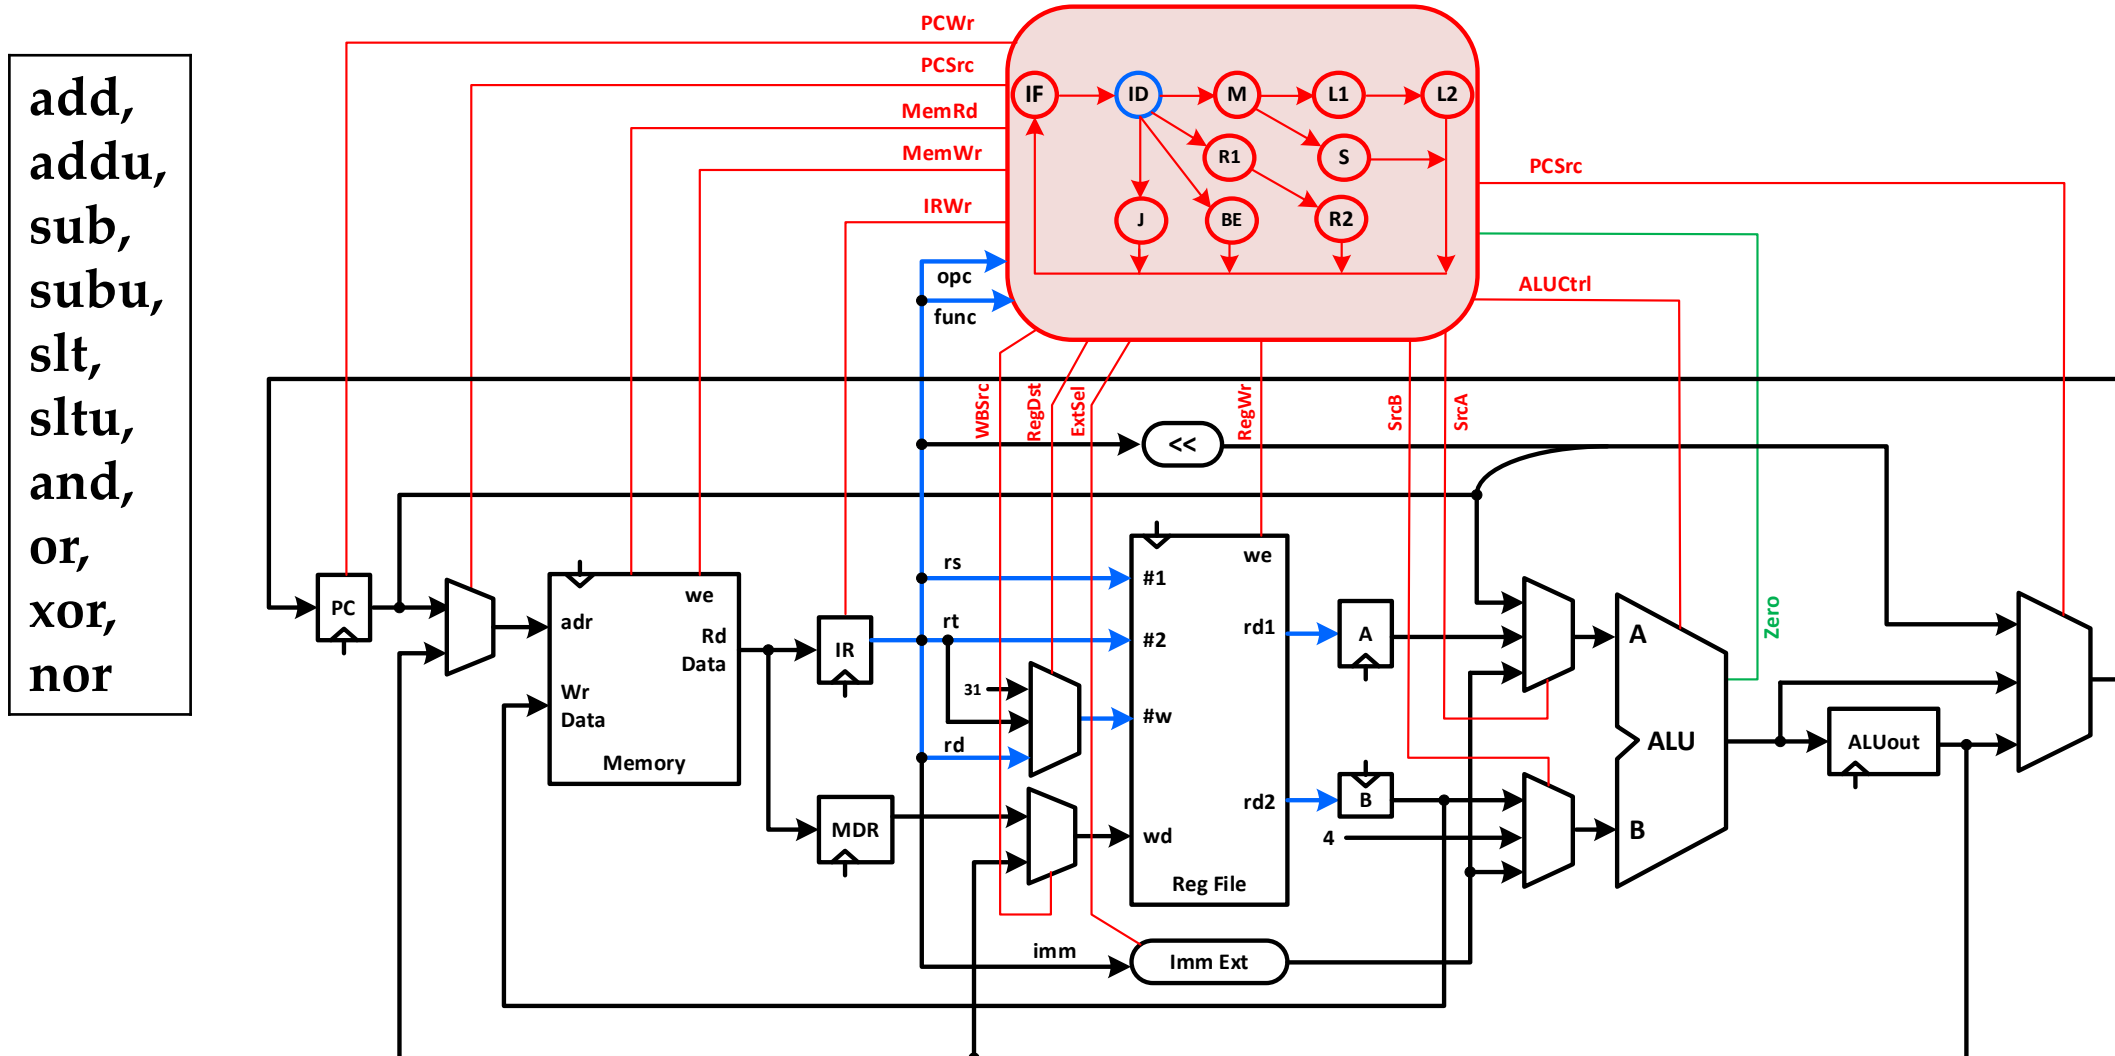
MIPS – Multi-Cycle Implementation



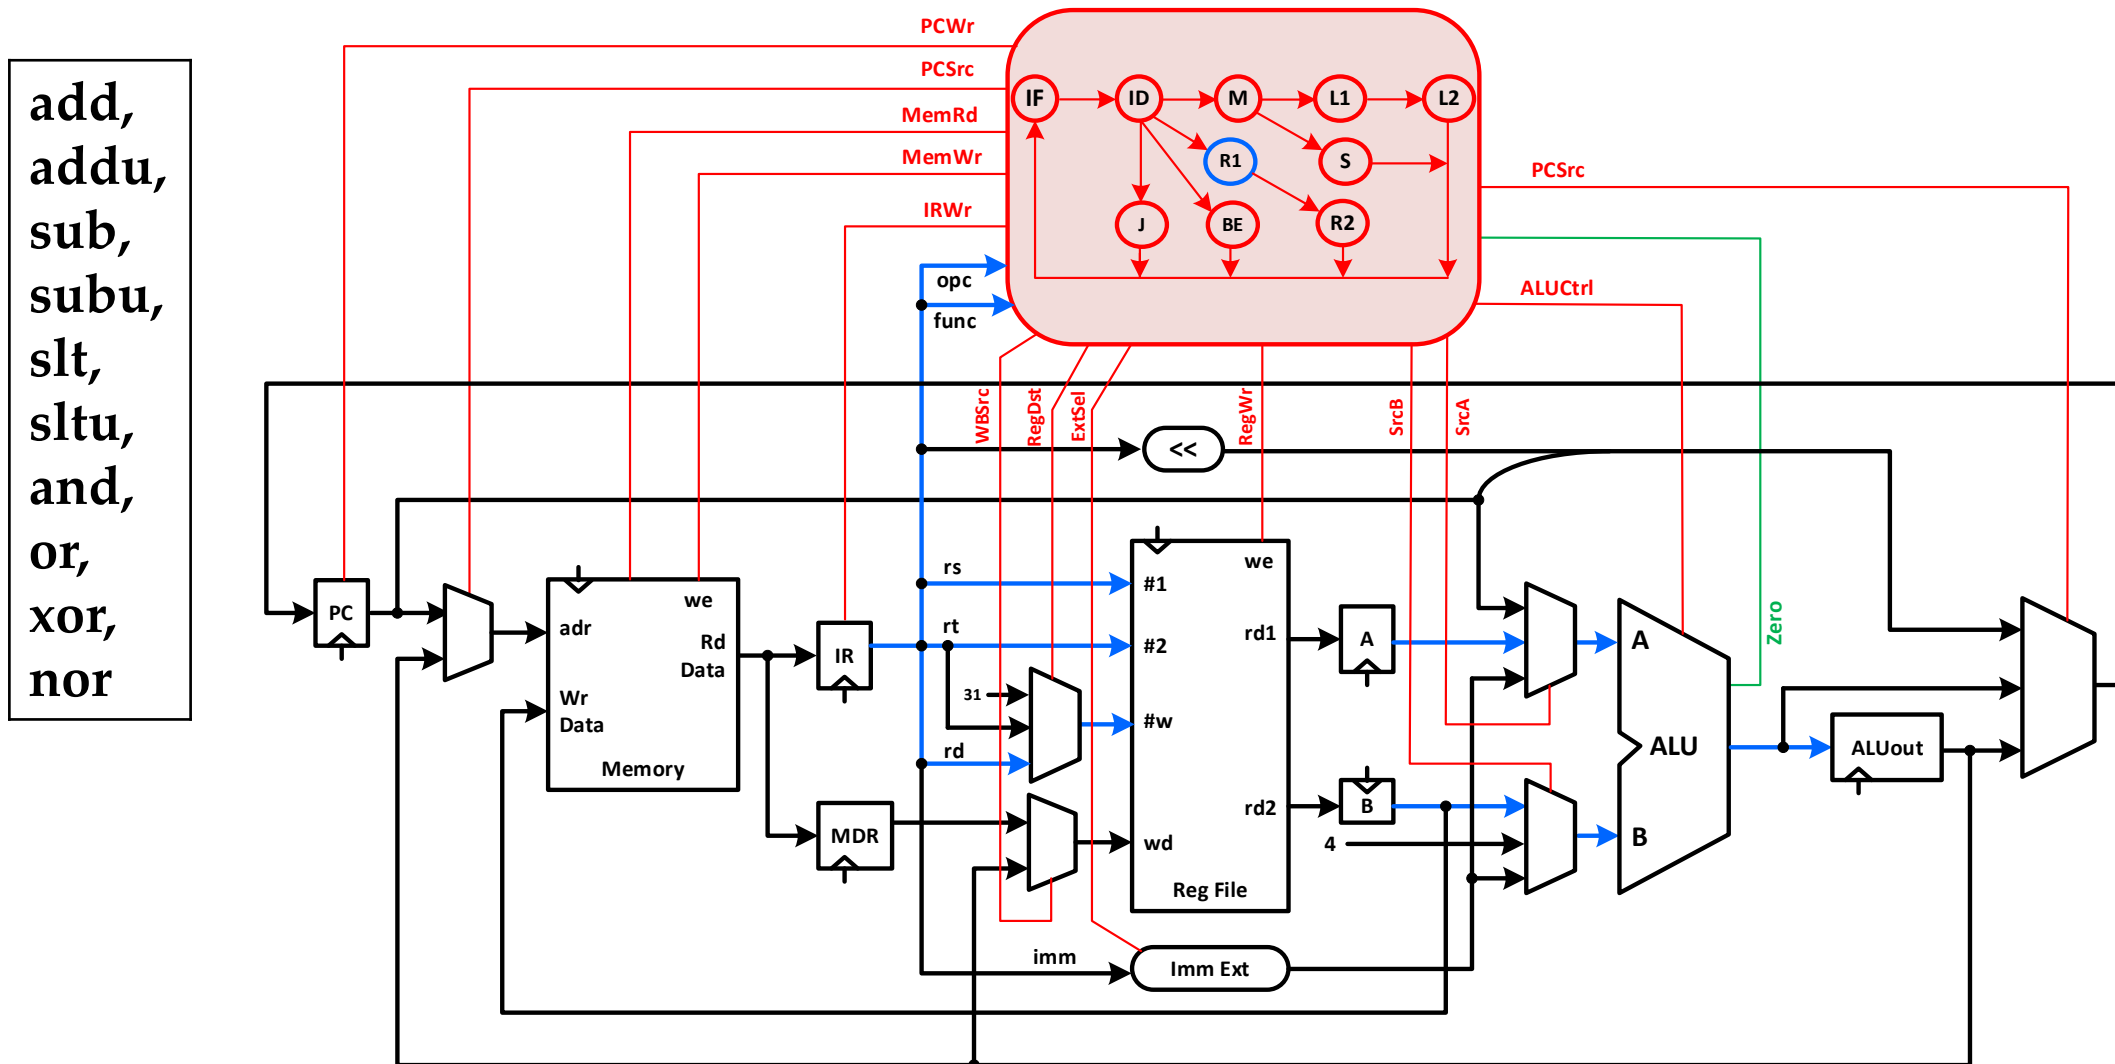
Arithmetic–Logic Instructions (Clk#1)



Arithmetic-Logic Instructions (Clk#2)

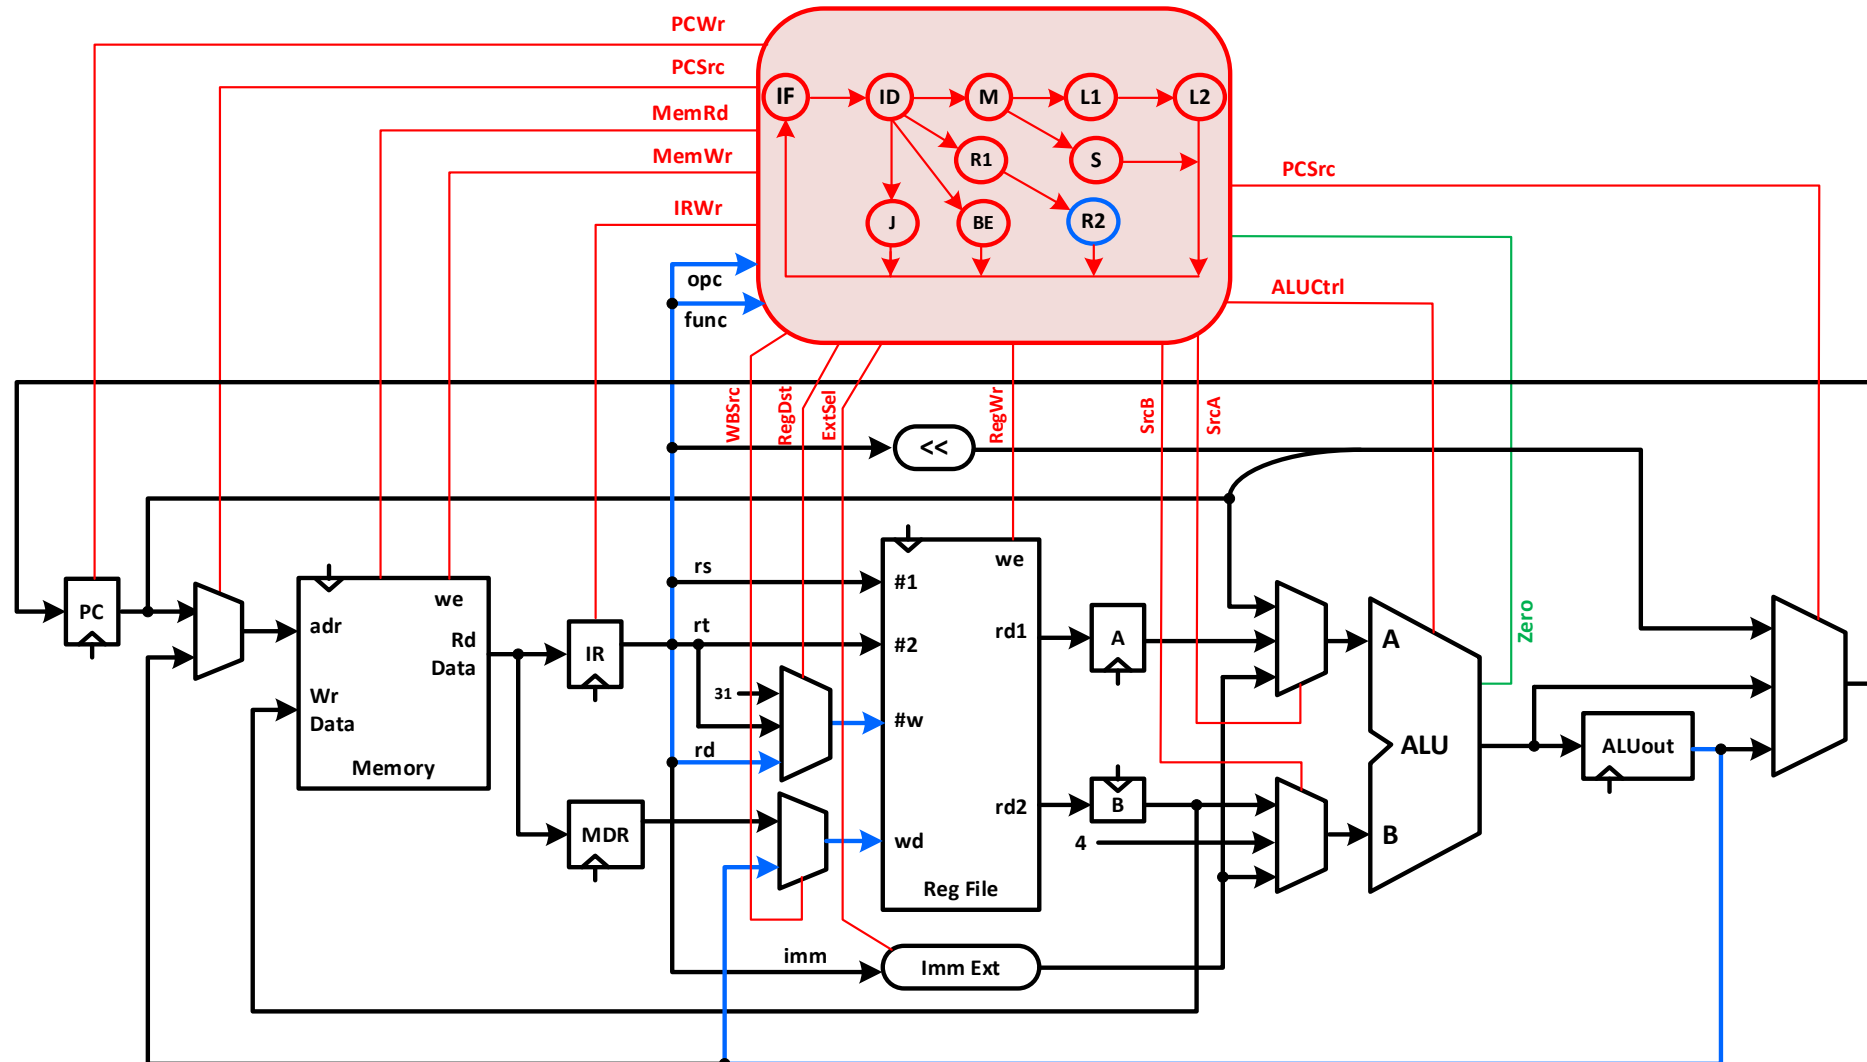


Arithmetic-Logic Instructions (Clk#3)



Arithmetic-Logic Instructions (Clk#4)

add,
addu,
sub,
subu,
slt,
sltu,
and,
or,
xor,
nor

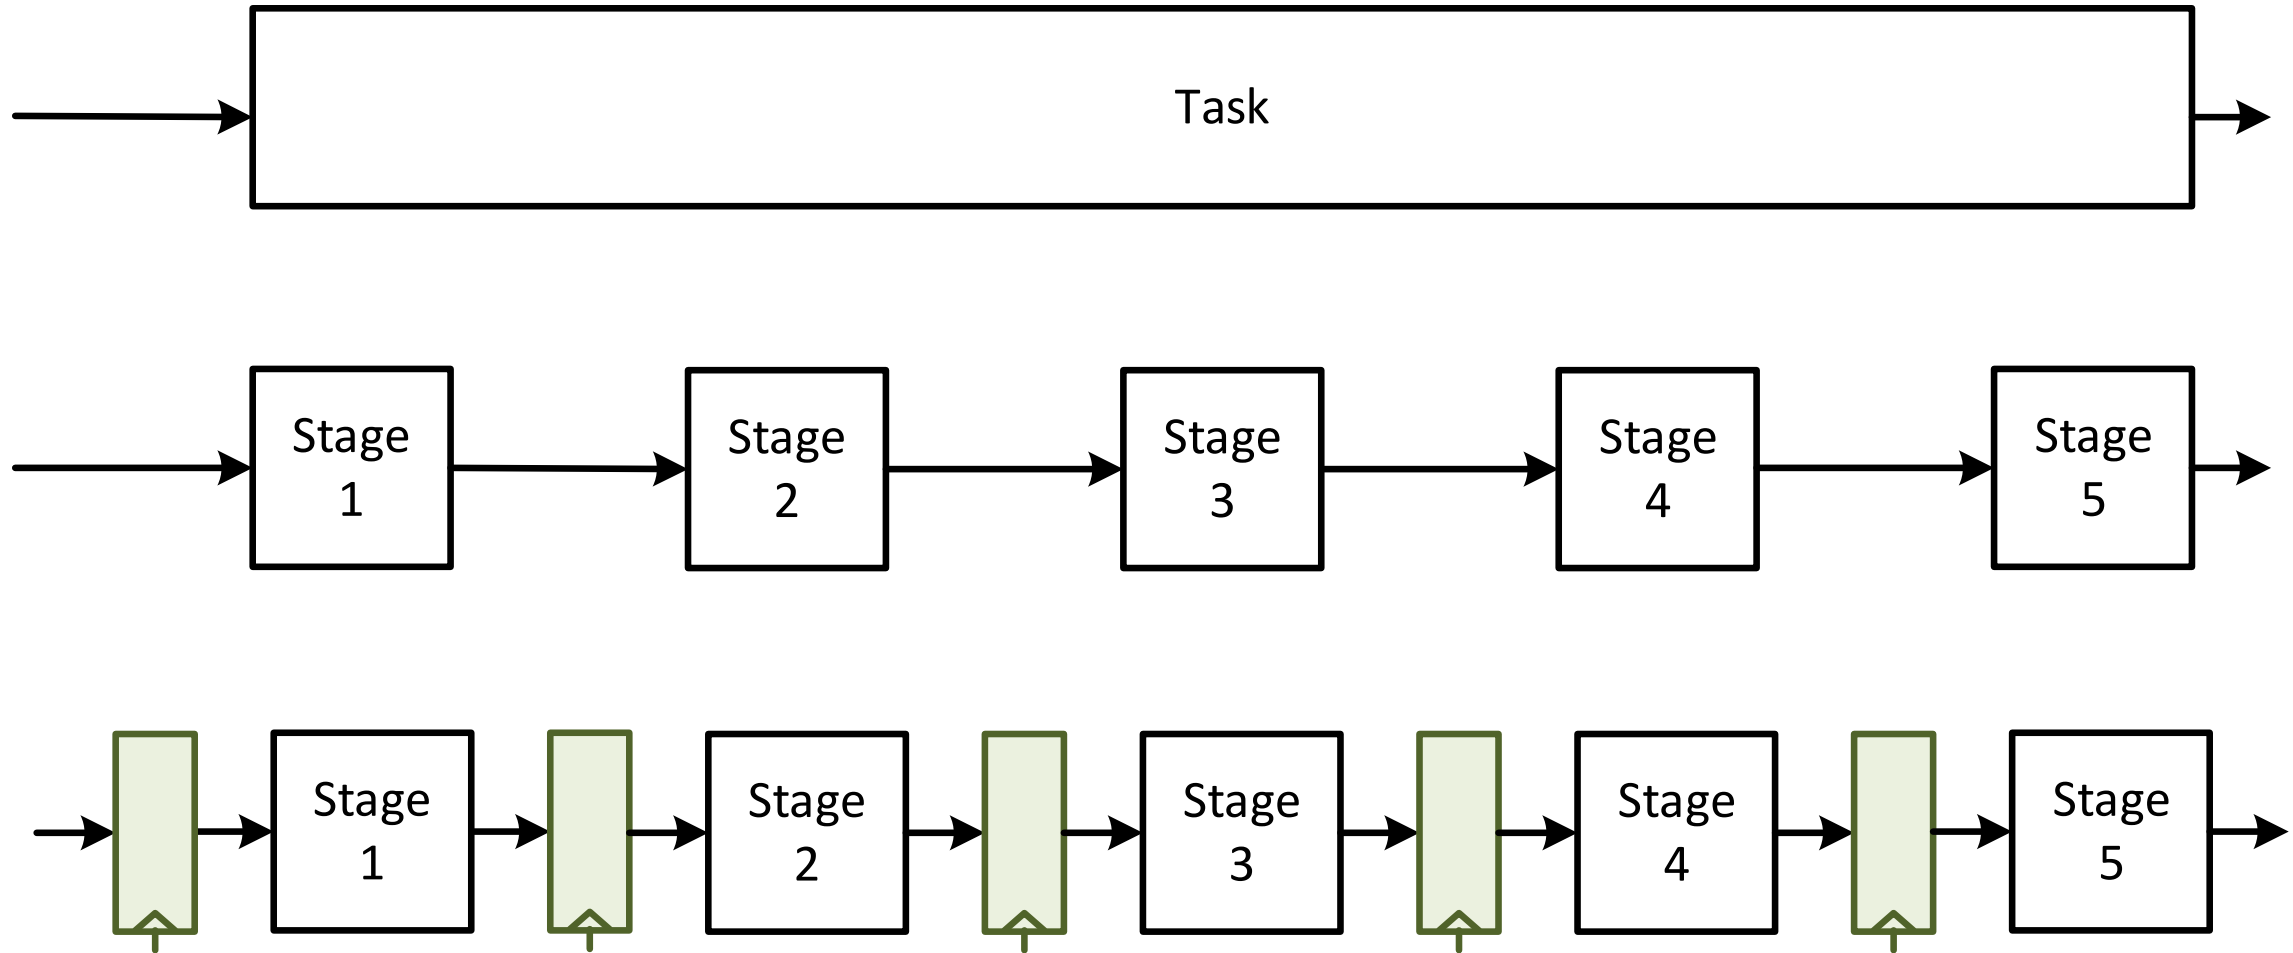


MIPS – Multi-Cycle Implementation

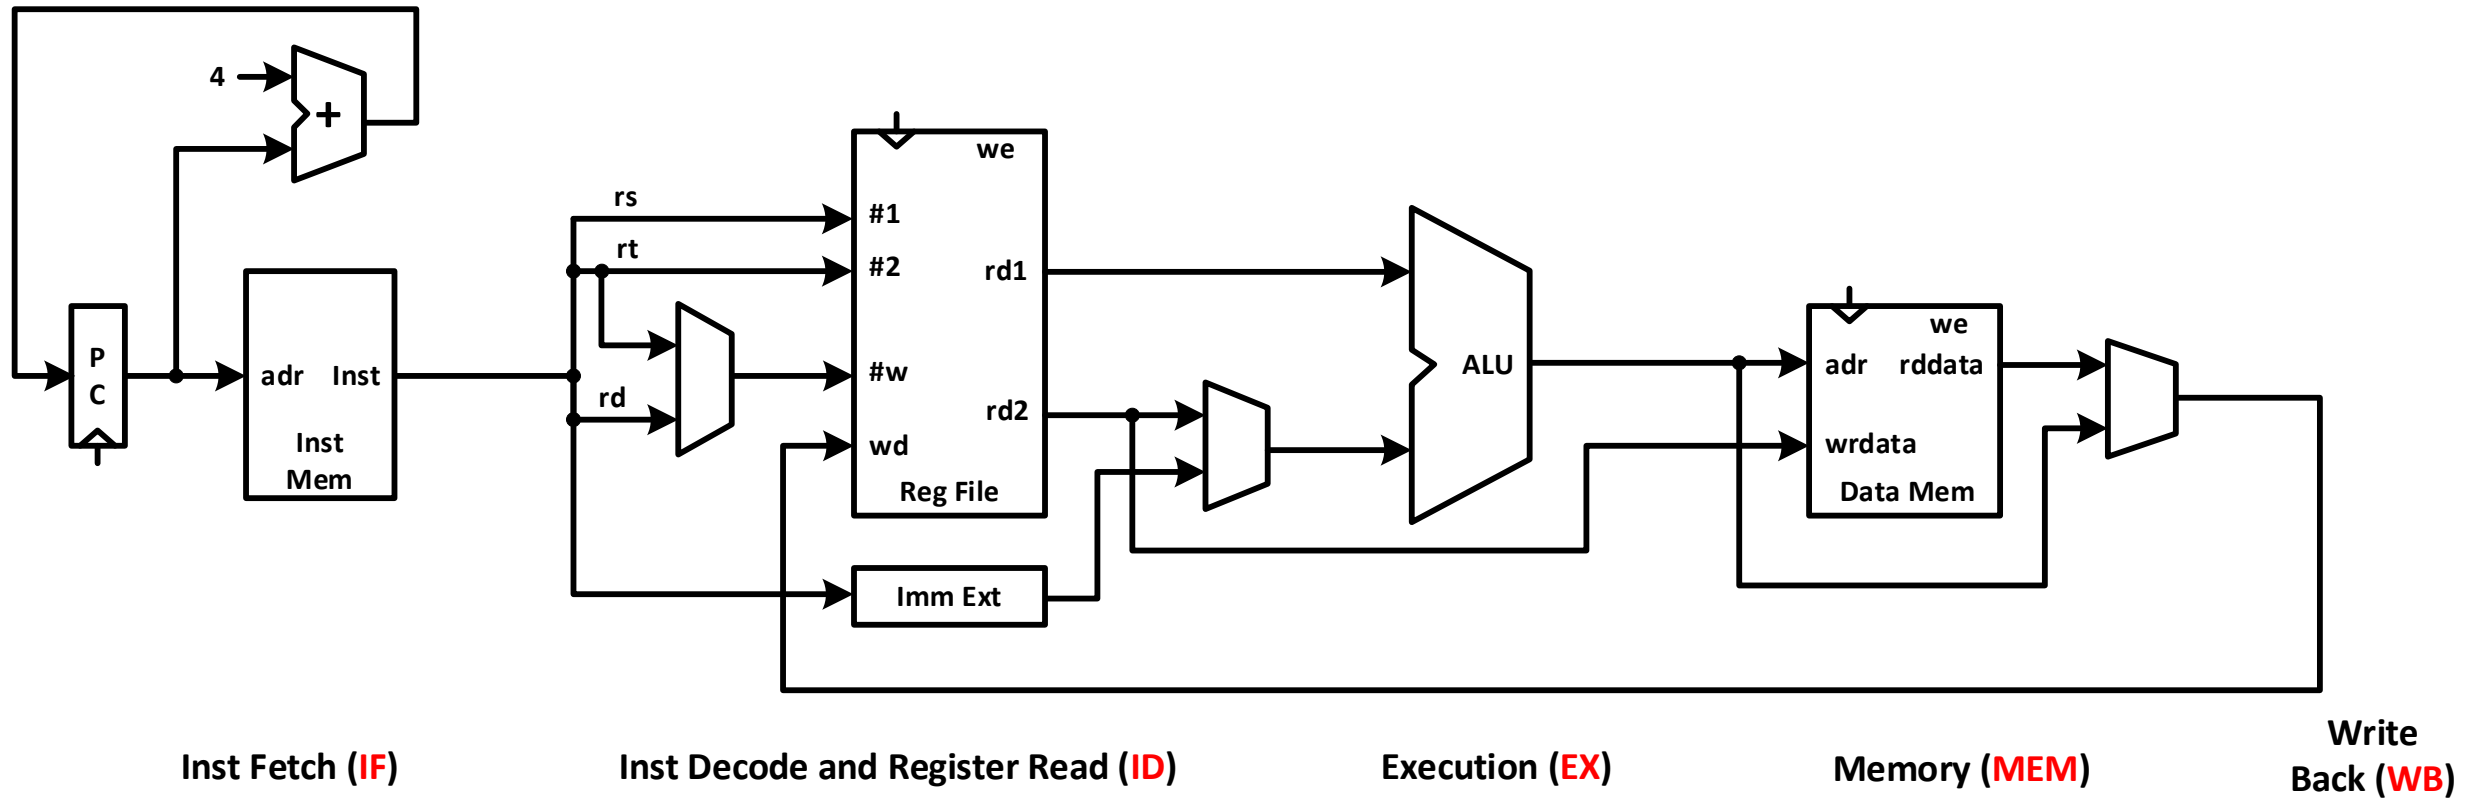
- Key features:
 - Data-path:
 - One ALU plays the role of ALU and two adders
 - One memory module plays the role of both data and instruction memory
 - Controller is implemented as a sequential circuit
 - Instructions need different number of clock cycles to execute (multi-cycle implementation)



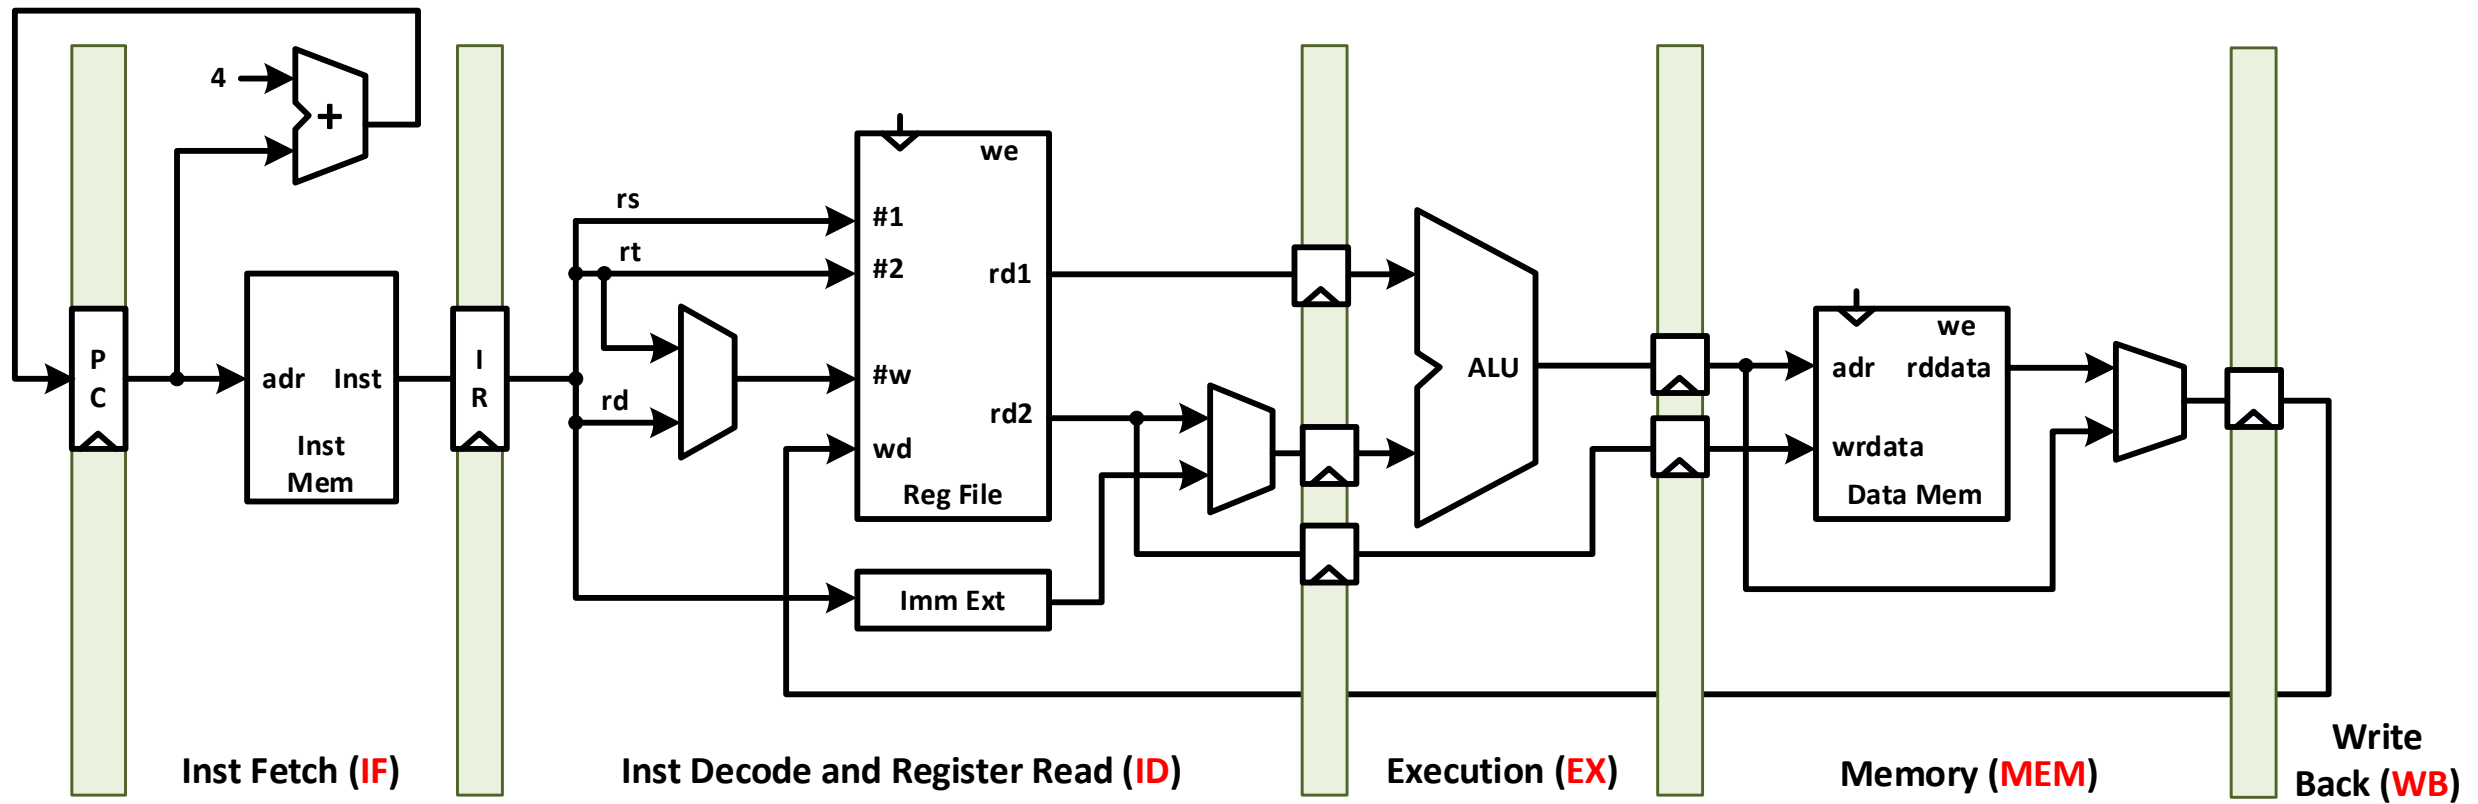
Pipelining – Idea



Simple Un-Pipelined Data-Path

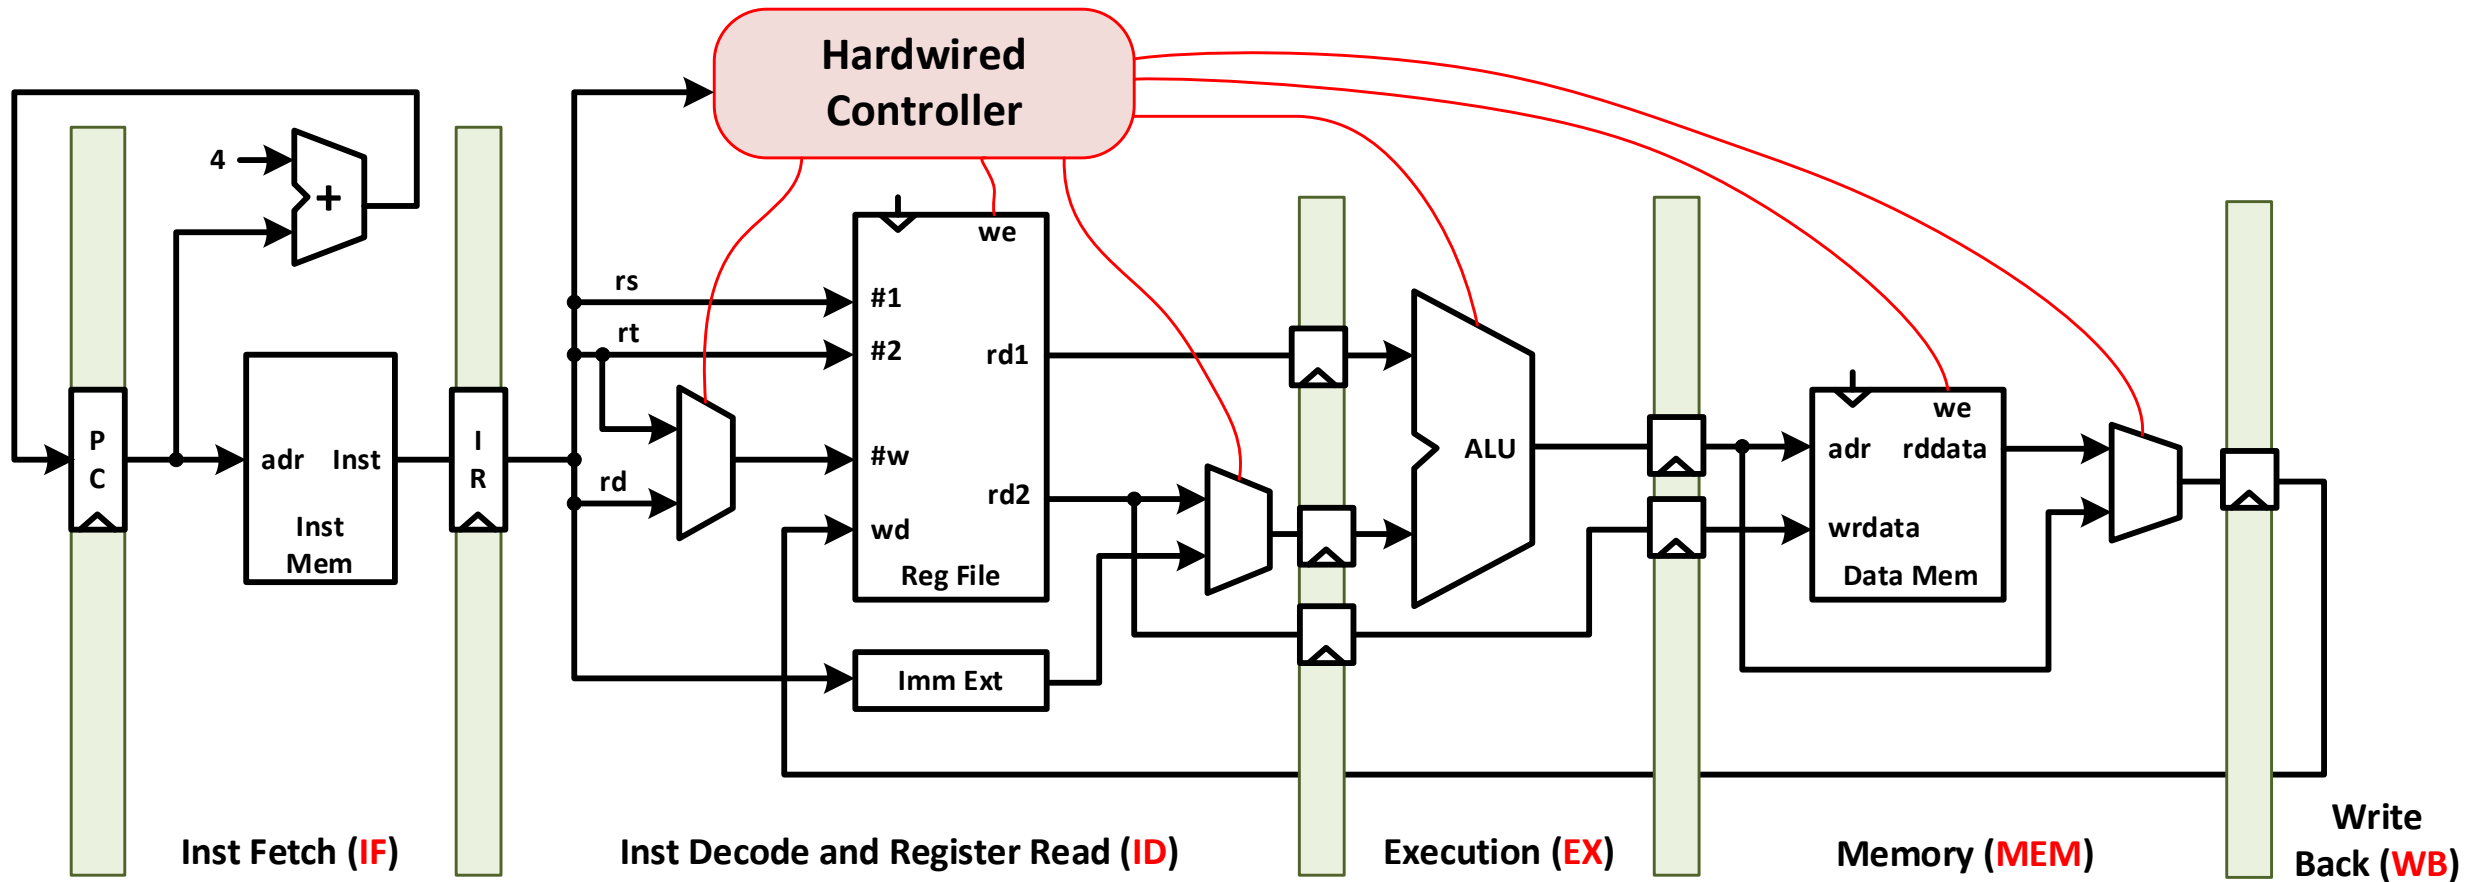


Pipelined Data-Path

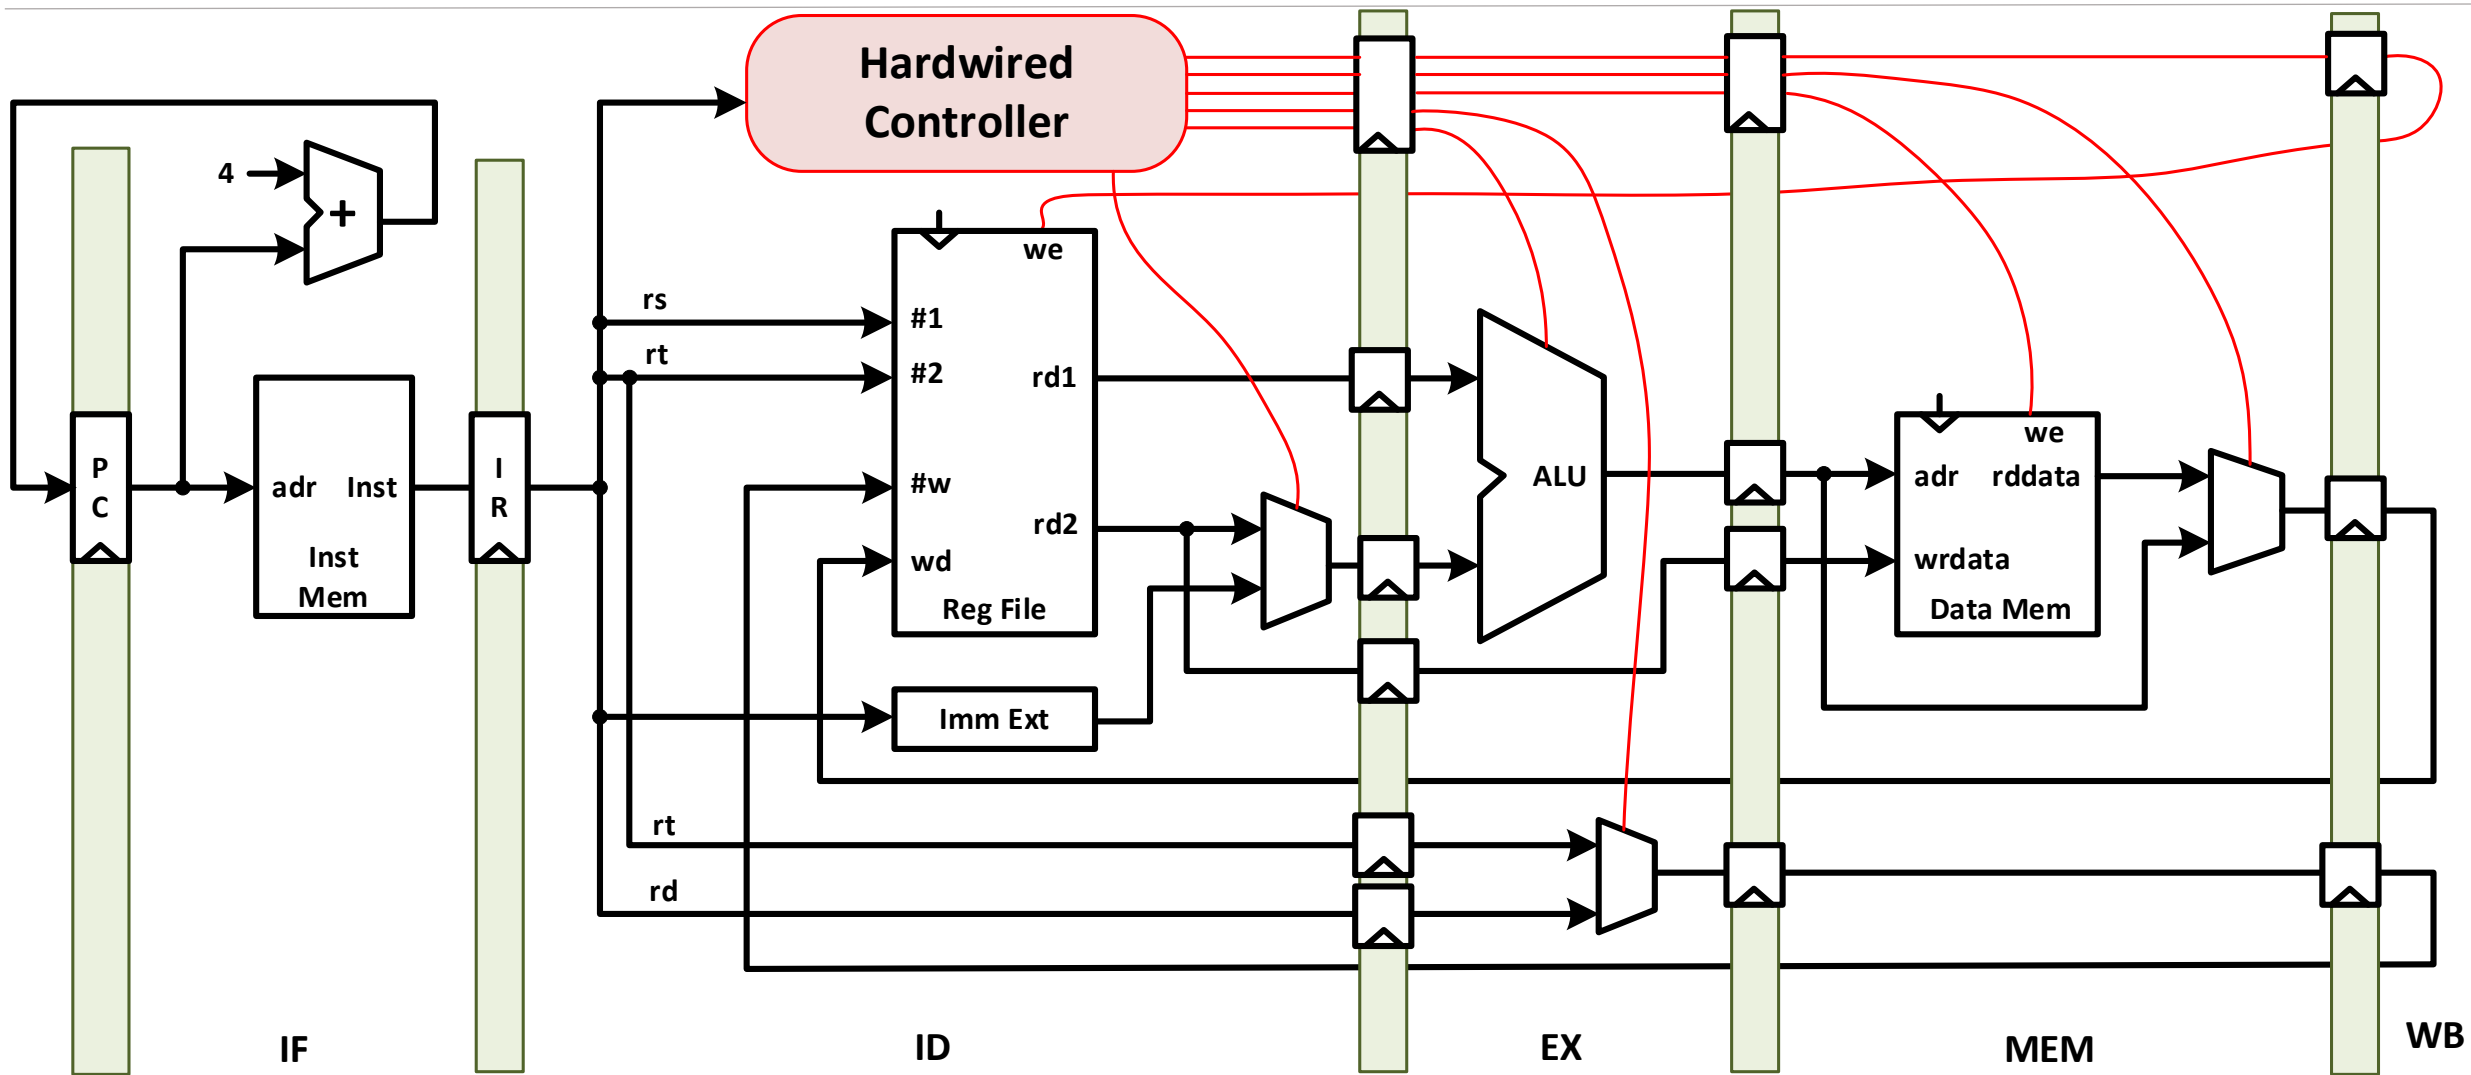


$$t_c > \max(t_{IM}, t_{RF}, t_{ALU}, t_{DM}, t_{WB}) \approx t_{DM}$$

Pipeline Controller



Pipeline Controller



Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Time}}{\text{Cycle}}$$

- **Inst/Prog:** Depends on source code, compiler technology, and ISA
- **Cycle/Inst (CPI):** Depends on ISA and micro-architecture
- **Time/Cycle:** Depends on micro-architecture and base technology

Processor Performance

| Micro-Architecture | CPI | Cycle Time |
|----------------------------|-------|------------|
| Multi-Cycle | > 1 | Short |
| Single Cycle (unpipelined) | 1 | Long |
| Pipelined | 1 | Short |

CPI Examples

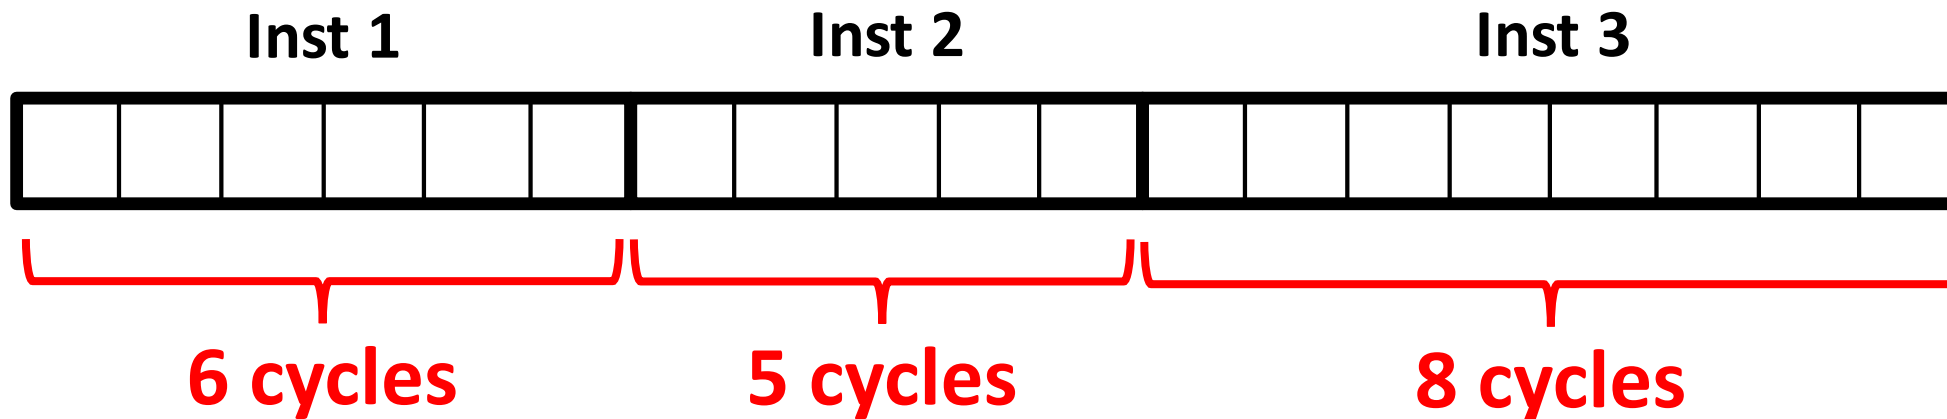
- Single-cycle implementation



- 3 instructions, 3 cycles, $\text{CPI} = 1$

CPI Examples

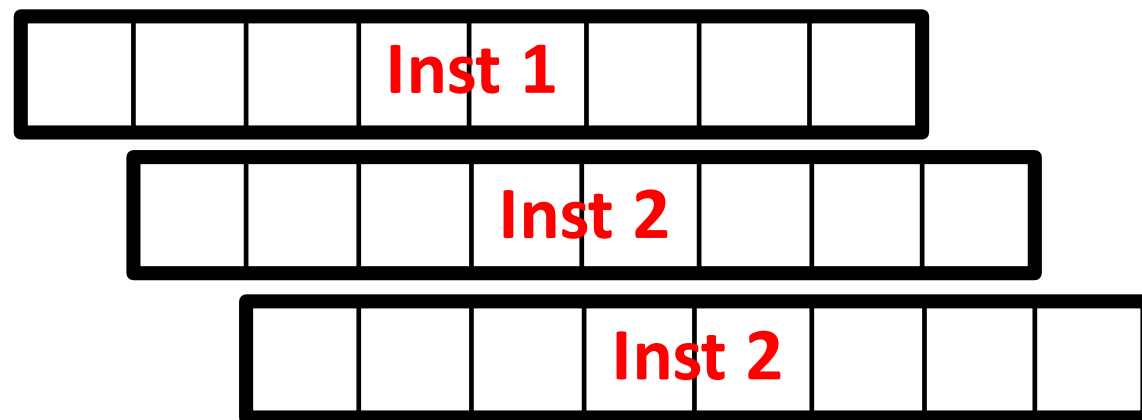
- Multi-cycle implementation



- 3 instructions, 19 cycles, $\text{CPI} = 6.33$

CPI Examples

- Pipeline implementation



- 3 instructions, 3 cycles, $CPI = 1$

Technology Assumption

- A small amount of very fast memory (Cache)
- Fast ALU
- Reasonable assumption:

$$t_{IM} \approx t_{RF} \approx t_{ALU} \approx t_{DM} \approx t_{WB}$$

- We focus on a 5-stage pipeline
 - Some commercial processors use 30-stage pipeline!



Technology Assumption

- Let's assume:

$$t_{IM} = t_{RF} = t_{ALU} = t_{DM} = t_{WB} = 200^{ps}$$

- Use a program with n instruction as a benchmark. Find the speedup of the pipelined implementation over SC implementation.

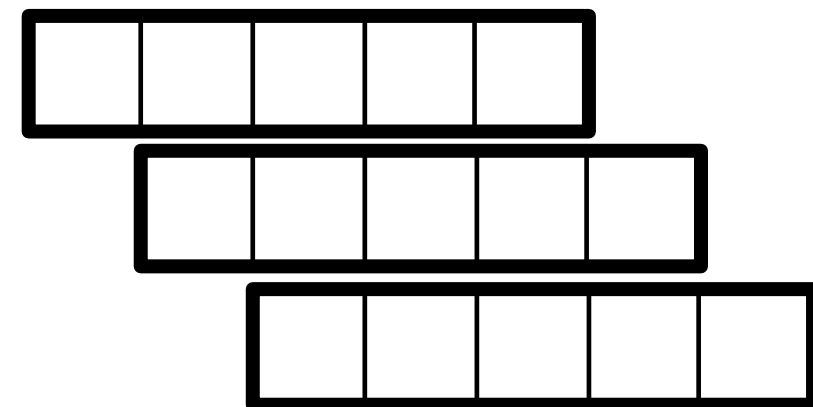
Technology Assumption

$$T_{SC} = (5 \times 200) \times n = 1000n$$

$$T_{Pipe} = (5 \times 200) + (n - 1) \times 200 = 800 + 200n$$

$$Speedup = \frac{T_{SC}}{T_{Pipe}} = \frac{1000n}{800 + 200n}$$

$$Speedup \xrightarrow{n \rightarrow +\infty} \frac{1000n}{200n} = 5$$



Generalization

- Let's assume a k -stage pipeline with the delays t_1, t_2, \dots, t_k , respectively
- Also assume: $t_1 + t_2 + \dots + t_k = T$
- Use a program with n instruction as a benchmark. Find the speedup of the pipelined implementation over SC implementation.

Generalization

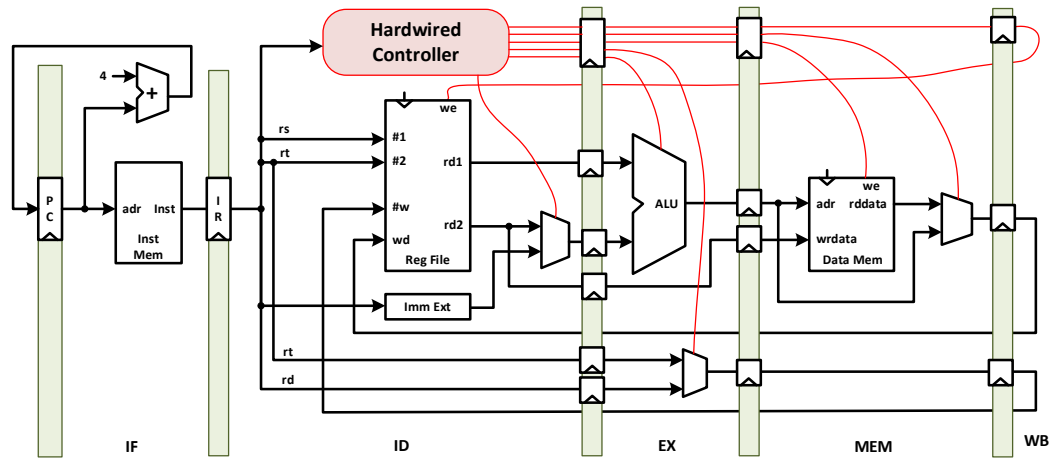
$$T_{SC} = T \times n$$

$$T_{Pipe} = k \times \text{Max}(t_1, \dots, t_k) + (n - 1) \times \text{Max}(t_1, \dots, t_k)$$

$$\text{Speedup} = \frac{T_{SC}}{T_{Pipe}} = \frac{T \times n}{k \times \text{Max}(t_1, \dots, t_k) + (n - 1) \times \text{Max}(t_1, \dots, t_k)}$$

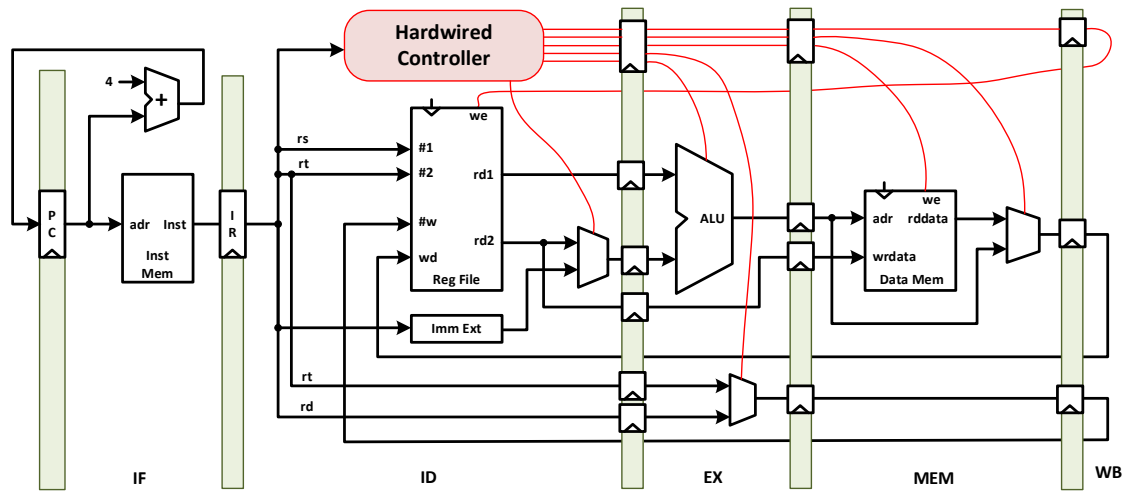
$$\text{Speedup} \xrightarrow{n \rightarrow +\infty} \frac{T}{\text{Max}(t_1, \dots, t_k)}$$

Pipeline Diagrams: Instructions vs. Time



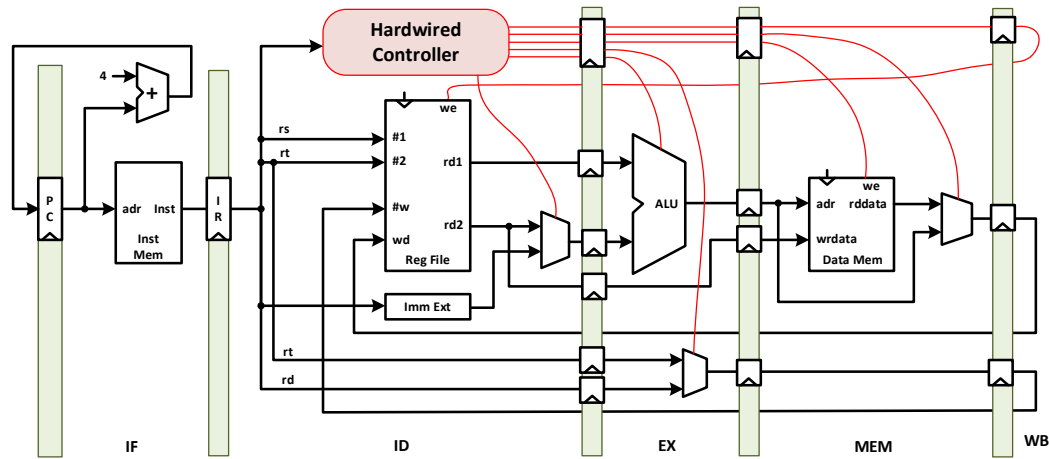
| Inst | C1 | C2 | C3 | C4 | C5 | C6 | C7 | C8 | ... |
|-------|-----------------|-----------------|-----------------|------------------|------------------|------------------|------------------|------------------|-----------------|
| Inst1 | IF ₁ | ID ₁ | EX ₁ | MEM ₁ | WB ₁ | | | | |
| Inst2 | | IF ₂ | ID ₂ | EX ₂ | MEM ₂ | WB ₂ | | | |
| Inst3 | | | IF ₃ | ID ₃ | EX ₃ | MEM ₃ | WB ₃ | | |
| Inst4 | | | | IF ₄ | ID ₄ | EX ₄ | MEM ₄ | WB ₄ | |
| Inst5 | | | | | IF ₅ | ID ₅ | EX ₅ | MEM ₅ | WB ₅ |

Pipeline Diagrams: Space vs. Time



| Time | C0 | C1 | C2 | C3 | C4 | C5 | C6 | C7 | ... |
|------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|
| IF | Inst ₁ | Inst ₂ | Inst ₃ | Inst ₄ | Inst ₅ | | | | |
| ID | | Inst ₁ | Inst ₂ | Inst ₃ | Inst ₄ | Inst ₅ | | | |
| EX | | | Inst ₁ | Inst ₂ | Inst ₃ | Inst ₄ | Inst ₅ | | |
| MEM | | | | Inst ₁ | Inst ₂ | Inst ₃ | Inst ₄ | Inst ₅ | |
| WB | | | | | Inst ₁ | Inst ₂ | Inst ₃ | Inst ₄ | Inst ₅ |

Pipeline Diagrams: Space vs. Time



| Inst | C1 | C2 | C3 | C4 | C5 | C6 | C7 | C8 | ... |
|------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|
| ID | Inst ₁ | Inst ₂ | Inst ₃ | Inst ₄ | Inst ₅ | | | | |
| ID | | Inst ₁ | Inst ₂ | Inst ₃ | Inst ₄ | Inst ₅ | | | |
| EX | | | Inst ₁ | Inst ₂ | Inst ₃ | Inst ₄ | Inst ₅ | | |
| MEM | | | | Inst ₁ | Inst ₂ | Inst ₃ | Inst ₄ | Inst ₅ | |
| WB | | | | | Inst ₁ | Inst ₂ | Inst ₃ | Inst ₄ | Inst ₅ |

Hazards

- **Structural hazard:** An instruction in the pipeline needs a resource being used by another instruction in the pipeline
- **Data hazard:** An instruction depends on a data value produced by an earlier instruction in the pipeline
- **Control hazard:** An instruction depends on a control decision made by an earlier instruction in the pipeline

Structural Hazard

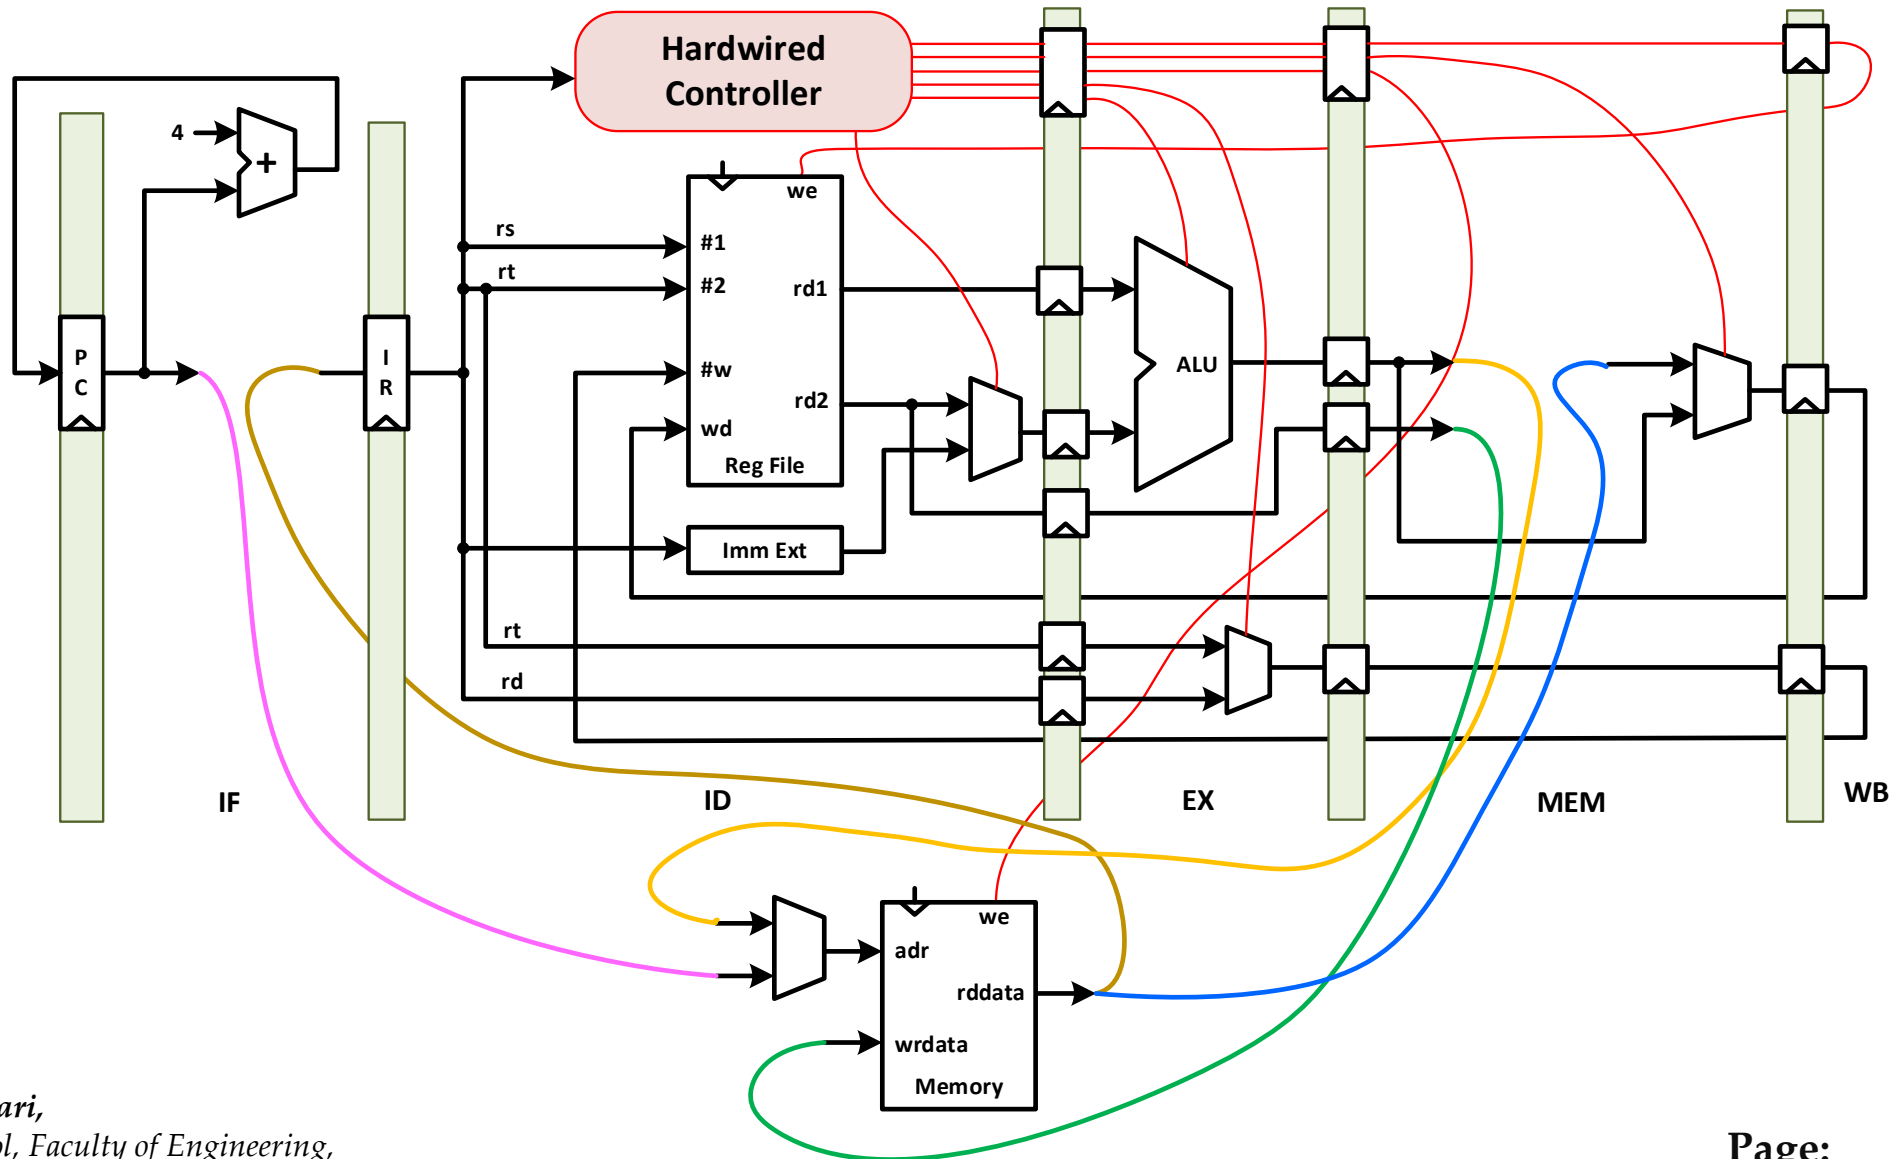
- **Schedule:** Programmer explicitly avoids scheduling instructions that would create structural hazards
- **Stall:** Hardware includes control logic that stalls until the earlier instruction is no longer using contended resource
- **Duplicate:** Add more hardware so that each instruction can access independent resources at the same time

Structural Hazard

- Simple 5-stage MIPS pipeline has no structural hazards specifically because ISA was designed that way



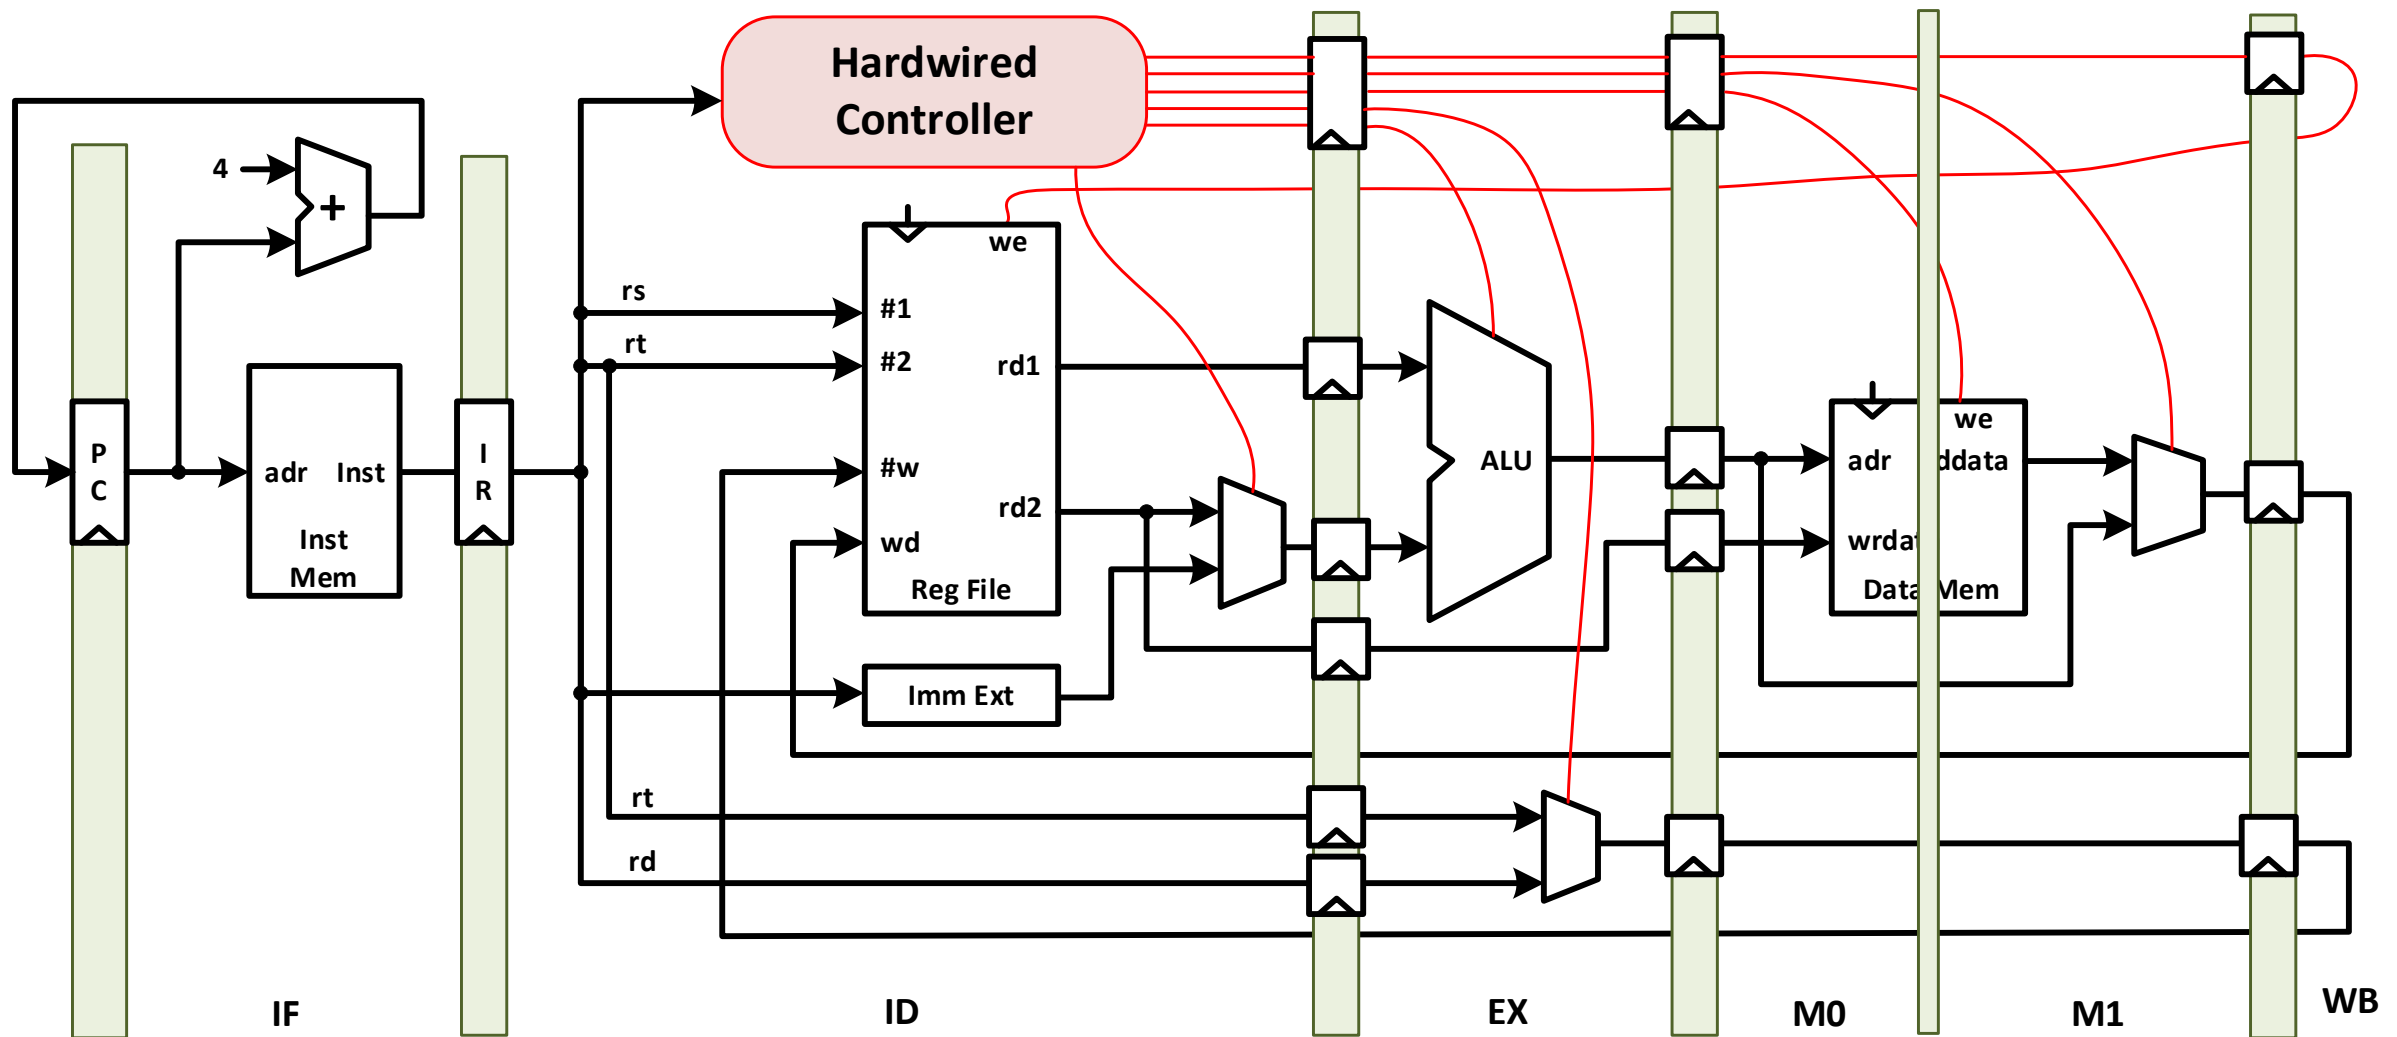
Example 1: Unified Memory



Example 1: Unified Memory

| Inst | C1 | C2 | C3 | C4 | C5 | C6 | C7 | C8 | C9 | C10 | C11 | C12 |
|------|----|----|----|----|----|----|----|----|----|-----|-----|-----|
| LW | F | D | X | M | W | | | | | | | |
| ADD | | F | D | X | M | W | | | | | | |
| ADD | | | F | D | X | M | W | | | | | |
| ADD | | | | – | F | D | X | M | W | | | |

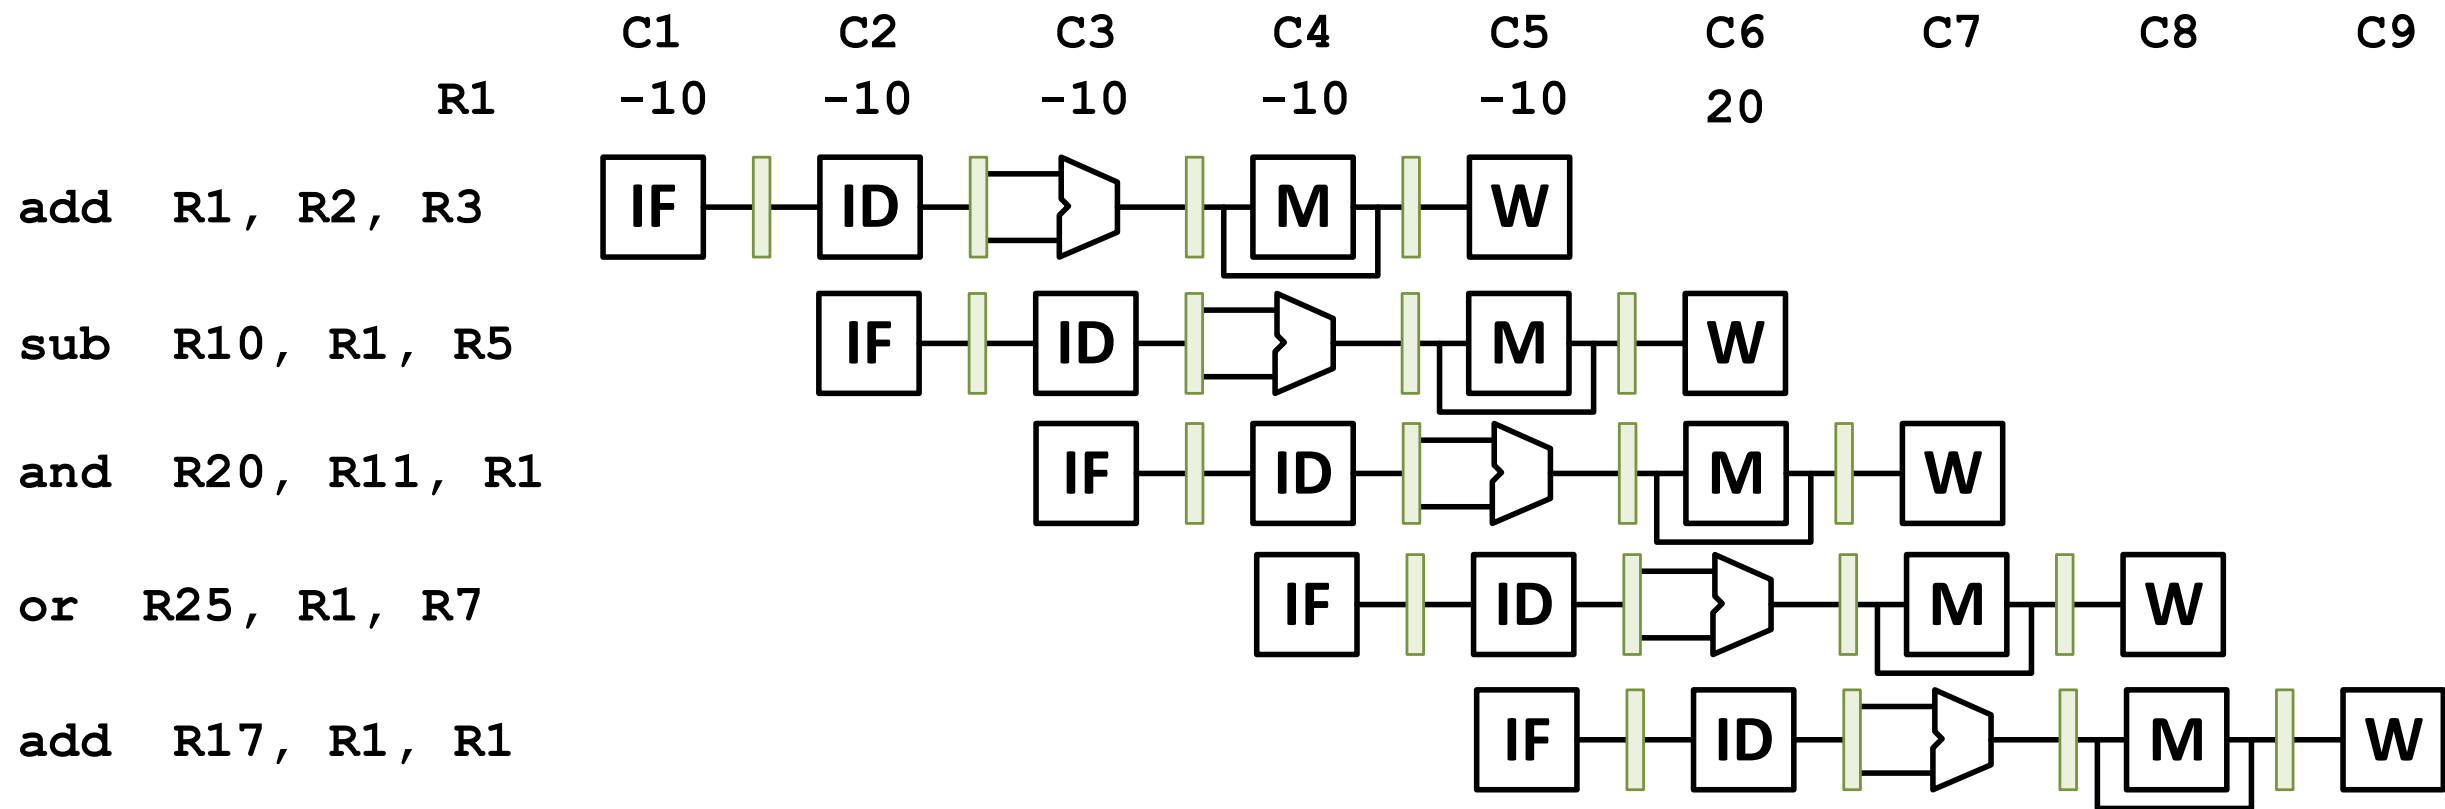
Example 2: 2-Cycle Memory



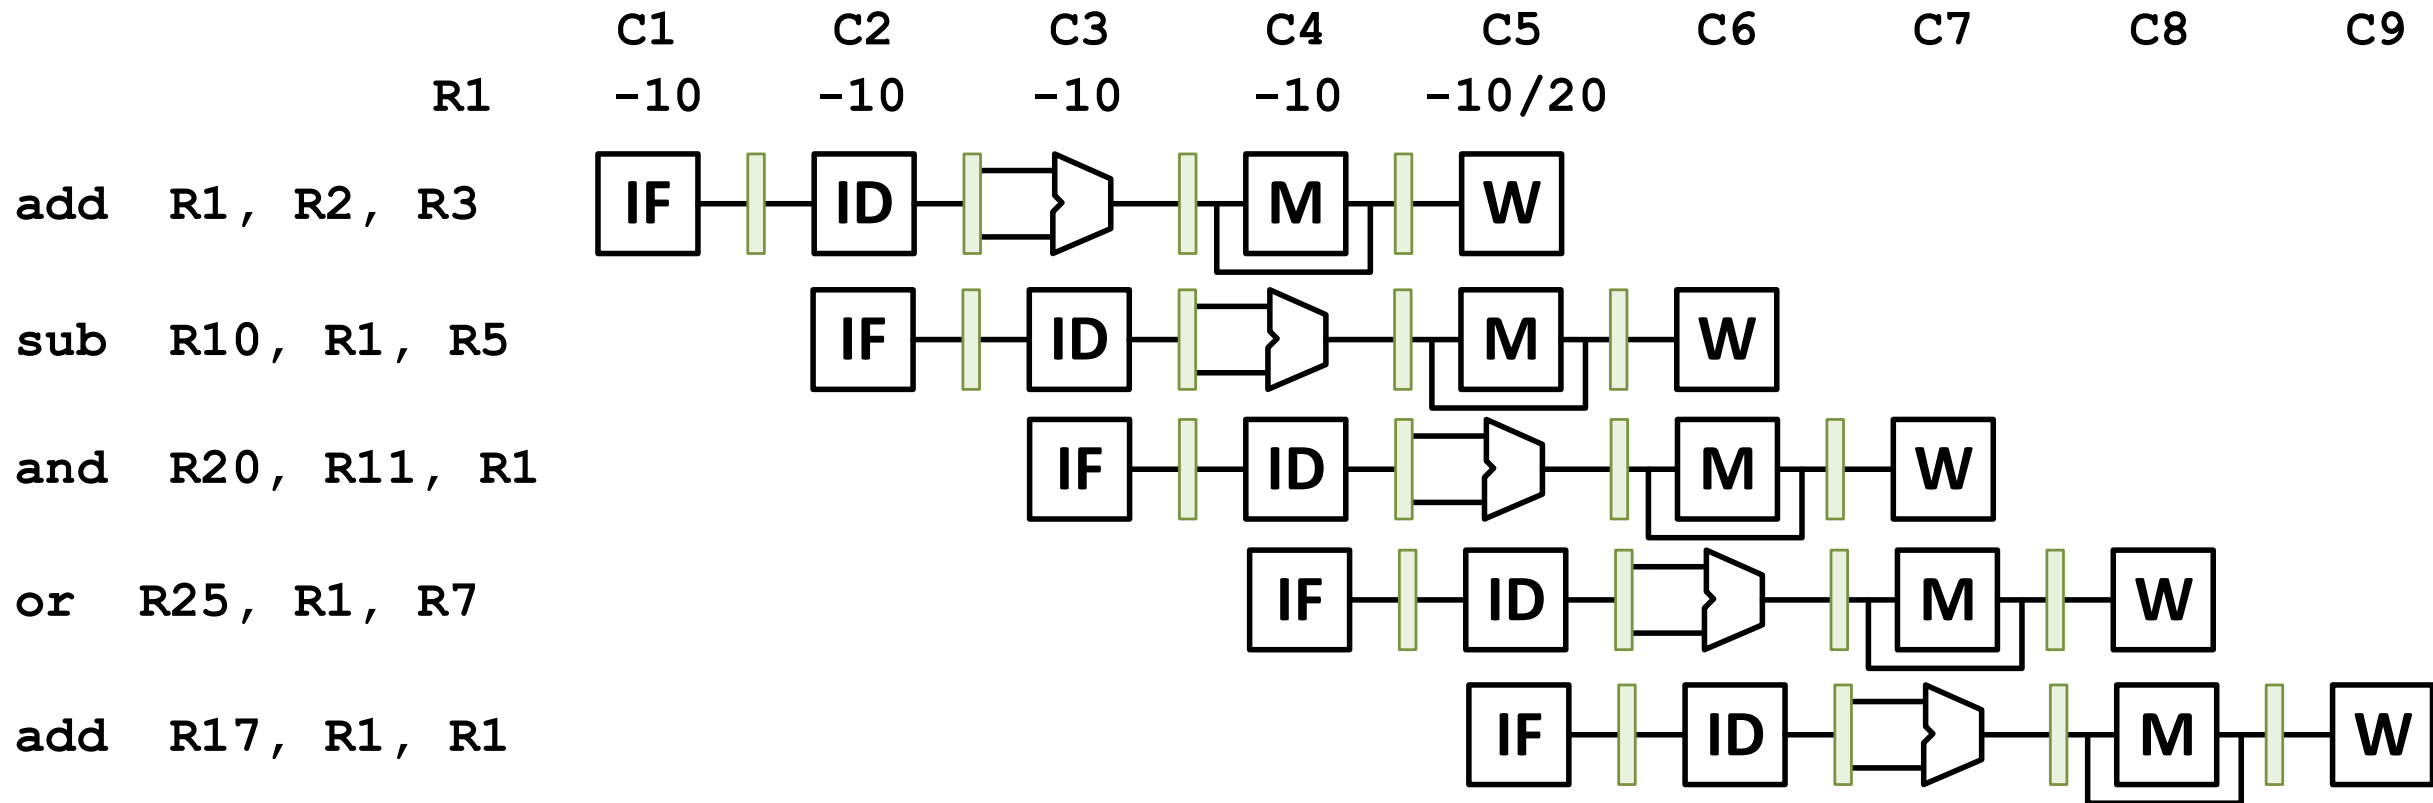
Example 2: 2-Cycle Memory

| Inst | C1 | C2 | C3 | C4 | C5 | C6 | C7 | C8 | C9 | C10 | C11 | C12 |
|------|----|----|----|----|----|----|----|----|----|-----|-----|-----|
| ADD | F | D | X | M0 | M1 | W | | | | | | |
| ADD | | F | D | X | M0 | M1 | W | | | | | |
| LW | | | F | D | X | M0 | M1 | W | | | | |
| LW | | | | F | D | X | – | M0 | M1 | W | | |
| ADD | | | | | F | D | – | X | M0 | M1 | W | |

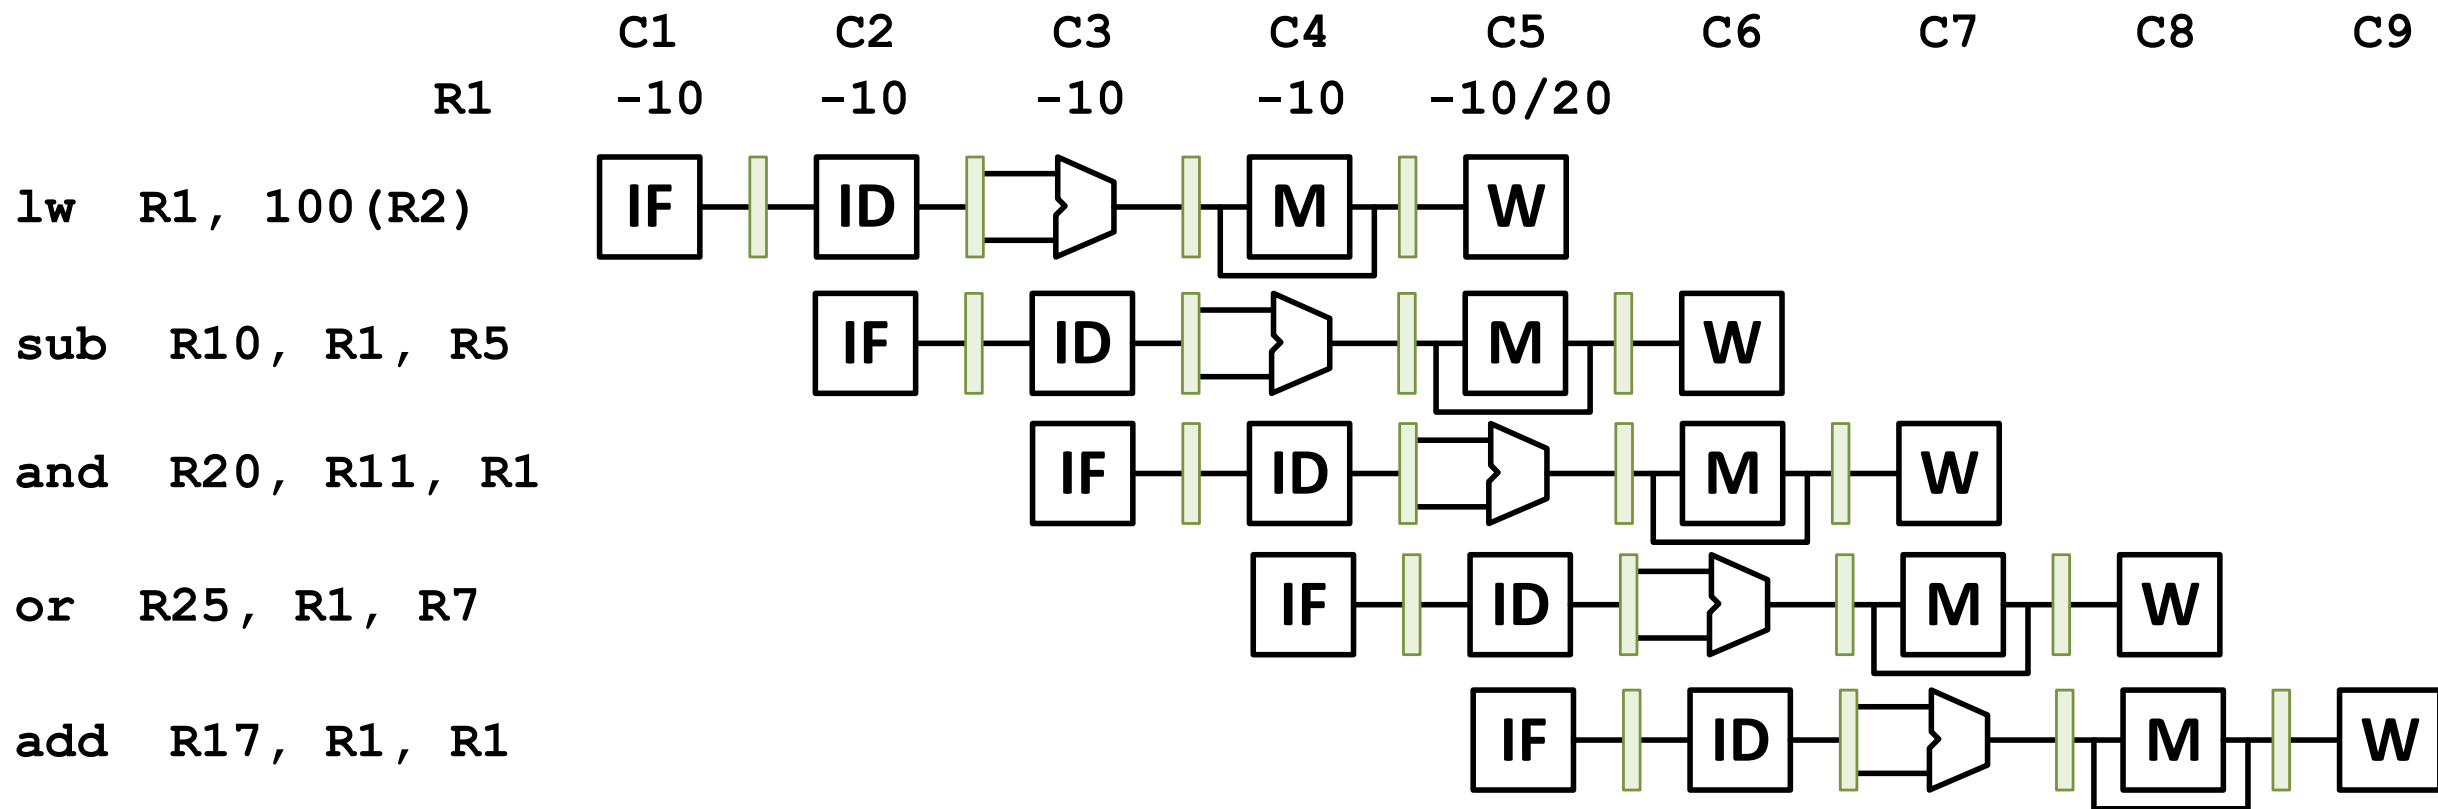
Data Hazard



Data Hazard



Data Hazard



Data Hazard

- **Schedule:** Programmer explicitly avoids scheduling instructions that would create data hazards
- **Stall:** Hardware includes control logic that freezes earlier stages until preceding instruction has finished producing data value
- **Bypass:** Hardware datapath allows values to be sent to an earlier stage before preceding instruction has left the pipeline
- **Speculate:** Guess that there is not a problem, if incorrect kill speculative instruction and restart